

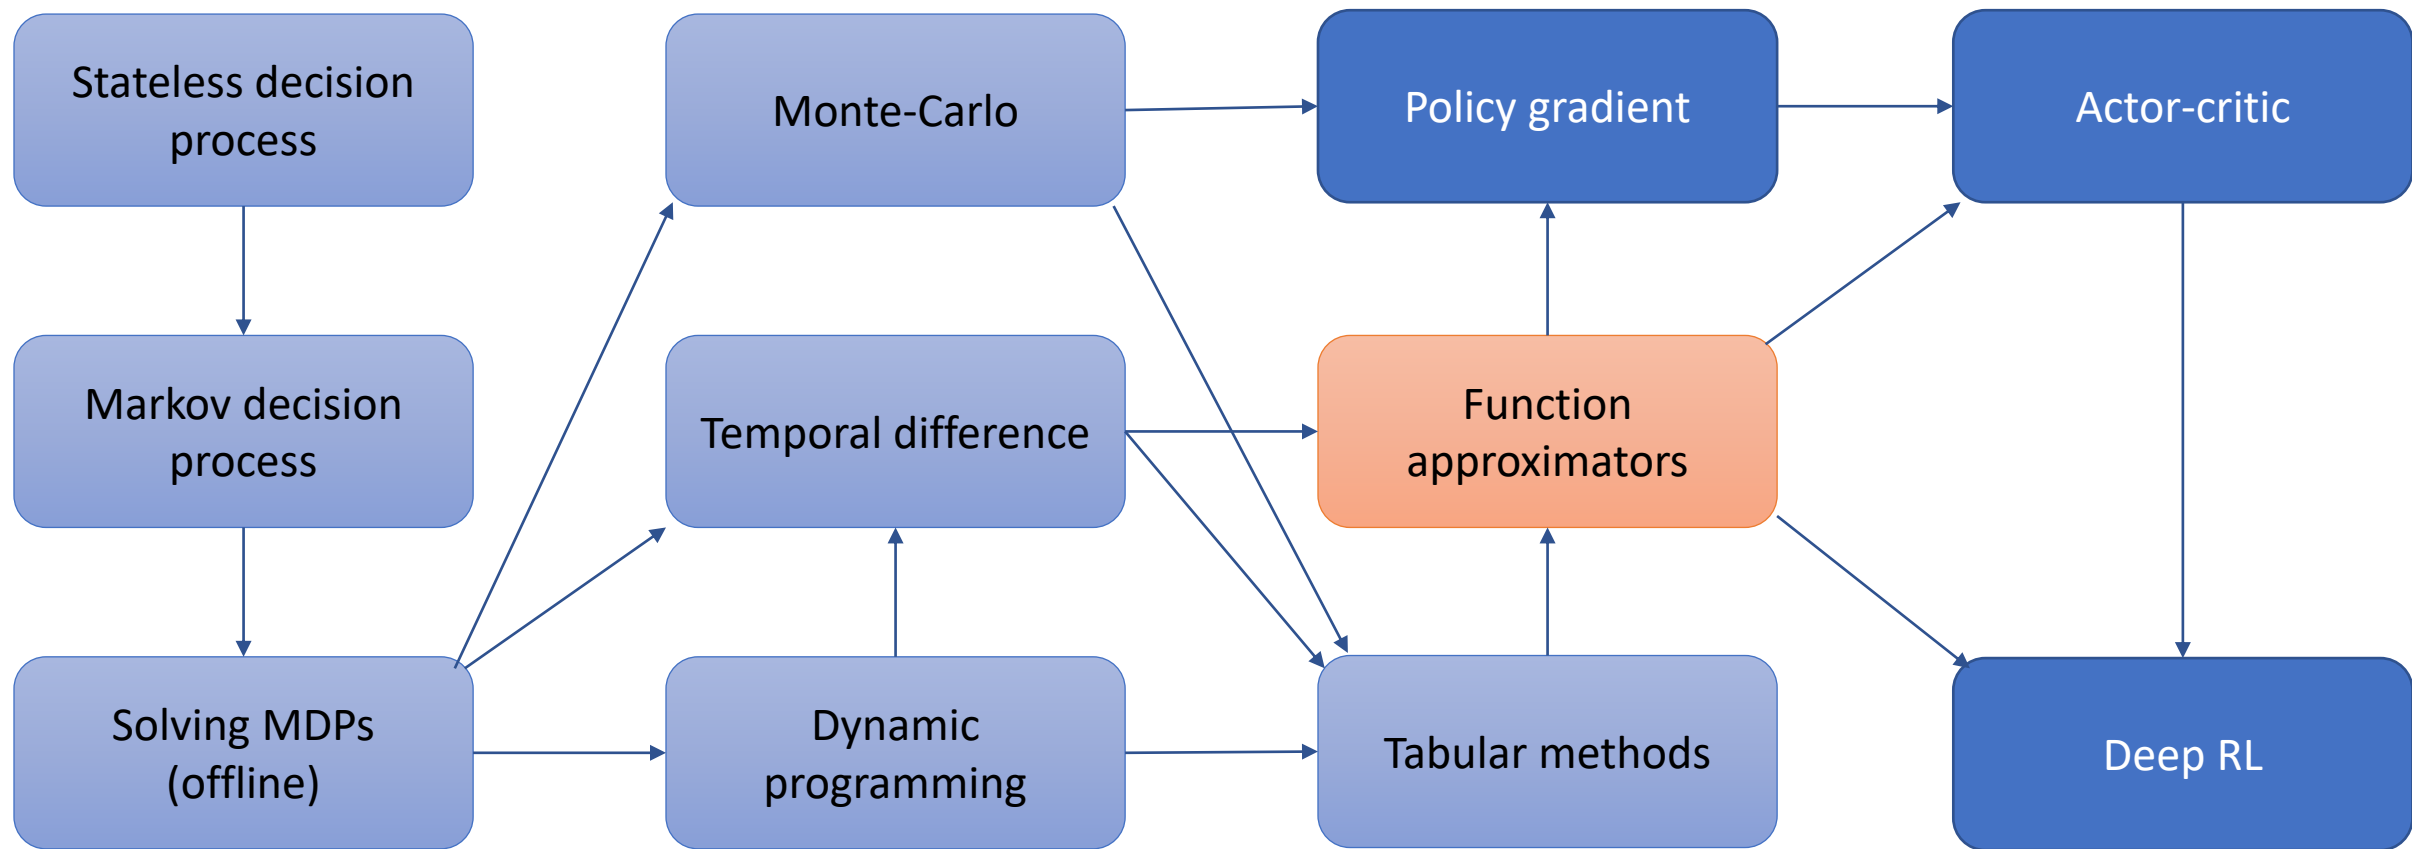
CSCE-642 Reinforcement Learning

Chapter 9: On-policy Prediction with Approximation





Instructor: Guni Sharon

CSCE-689, Reinforcement Learning




Generalizing value-based learning

- Tabular Learning keeps a table of all state values
- In realistic situations, we cannot possibly learn about every single state!
 - Too many states to visit them all during training
 - Too many states to hold a value table in memory
- Instead, we want to generalize:
 - Train on a small number of states from experience
 - Generalize that experience to new, similar situations
 - This is a fundamental idea in machine learning




Example: Pacman


Let's say we discover through experience that this state is bad:



In naïve tabular-learning, we know nothing about this state:




Or even this one!



- Naïve Q-learning
- After 50 training episodes



Learn an approximation function




- Q-learning with function approximator



Parametrized function approximator

- Assume that each state is vector of features (f_1, f_2, \dots, f_n) , e.g.,
 - Packman_location, Ghost1_location, Ghost2_location, food_location
 - Or even screen pixels
- A parametrized value approximator $\hat{v}(s; \theta)$ might look like this:
 - $= \sum_i \theta_i f_i$ or this: $= \sum_i \theta_i \sin(i f_i)$ or even this:
- Assume we know the true value for a set of states:
 - $v(S_1) = 5, v(S_2) = 8, v(S_3) = 2$
 - How can we update θ to reflect this information?




Gradient Decent

- Given: $v(S_1) = 5$, $v(S_2) = 8$, $v(S_3) = 2$
- We want to set θ such that $\forall s, \hat{v}(s; \theta) = v(s)$
 - Not possible in the general case, why?
 - Instead we'll try to minimize the errors: $\text{loss} = \sum_s |v(s) - \hat{v}(s; \theta)|$
 - Partial derivative of the loss with respect to θ_i = how to change θ_i such that loss will increase the most
 - Go the other way -> decrease loss
 - Ooops! Absolute value is not differentiable -> can't compute gradients
 - Simple fix: $\text{loss} = \frac{1}{2} \sum_s [v(s) - \hat{v}(s; \theta)]^2$ = squared loss function

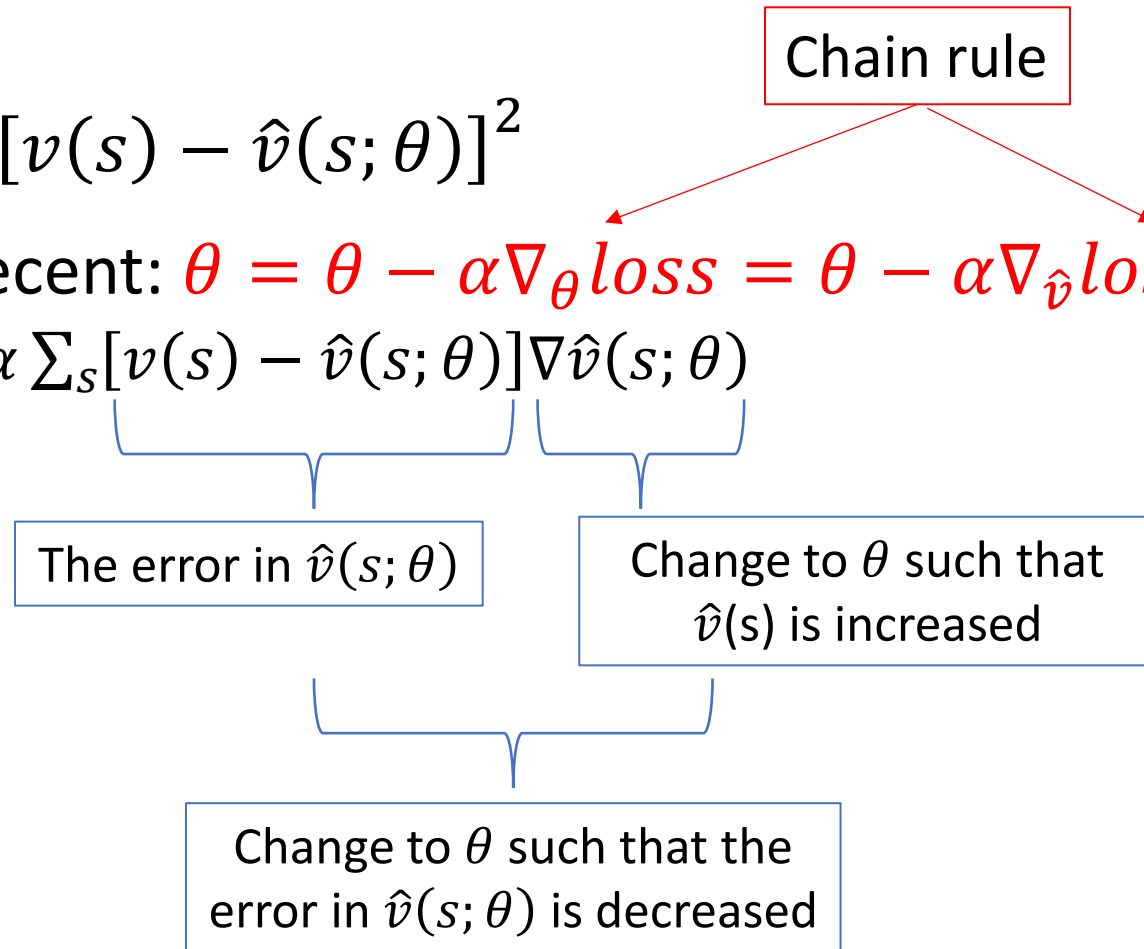
Gradient Decent

- $\text{loss} = \frac{1}{2} \sum_s [\nu(s) - \hat{\nu}(s; \theta)]^2$
- For each i
 - Push θ_i towards a direction that minimizes loss
 - $\theta_i = \theta_i - \frac{\partial \text{loss}}{\partial \theta_i}$
- More generally $\theta = \theta - \alpha \nabla \text{loss}$
 - $\nabla \text{loss} = \left(\frac{\partial \text{loss}}{\partial \theta_1}, \frac{\partial \text{loss}}{\partial \theta_2}, \dots, \frac{\partial \text{loss}}{\partial \theta_n} \right)$
 - α is the learning rate, requires tuning per domain, too large causes divergence to small results in slow learning or even premature convergence




Gradient Decent

- $\text{loss} = \frac{1}{2} \sum_s [v(s) - \hat{v}(s; \theta)]^2$
- Gradient Decent: $\theta = \theta - \alpha \nabla_{\theta} \text{loss} = \theta - \alpha \nabla_{\hat{v}} \text{loss} \cdot \nabla_{\theta} \hat{v}$
 - $\theta = \theta + \alpha \sum_s [v(s) - \hat{v}(s; \theta)] \nabla \hat{v}(s; \theta)$



Gradient Descent

- Idea:
 - Start somewhere
 - Repeat: Take a step in the gradient direction



Batch Gradient Decent

Minimize squared loss: $l(\theta) = \frac{1}{2} \sum_s [v(s) - \hat{v}(s; \theta)]^2$

- `init`
- `for iter = 1, 2, ...`
 - $\theta = \theta + \alpha \sum_s [v(s) - \hat{v}(s; \theta)] \nabla \hat{v}(s; \theta)$

Stochastic Gradient Decent (SGD)

Minimize squared loss: $l(\theta) = \frac{1}{2} \sum_s [v(s) - \hat{v}(s; \theta)]^2$

Observation: once gradient on one training example has been computed, might as well incorporate before computing next one

```
• init  $\theta$   
• for iter = 1, 2, ...  
    • pick random  $j$   
    •  $\theta = \theta + \alpha [v(s_j) - \hat{v}(s_j; \theta)] \nabla \hat{v}(s_j; \theta)$ 
```

Mini-Batch Gradient Decent

Minimize squared loss: $l(\theta) = \frac{1}{2} \sum_s [v(s) - \hat{v}(s; \theta)]^2$

Observation: gradient over small set of training examples (=mini-batch) can be computed in parallel, might as well do that instead of a single one

- `init θ`
- `for iter = 1, 2, ...`
 - `pick random subset of training examples J`
 - `$\theta = \theta + \alpha \sum_{s \in J} [v(s) - \hat{v}(s; \theta)] \nabla \hat{v}(s; \theta)$`

SGD for Monte Carlo estimation

Gradient Monte Carlo Algorithm for Estimating $\hat{v} \approx v_\pi$

Input: the policy π to be evaluated

Input: a differentiable function $\hat{v} : \mathcal{S} \times \mathbb{R}^d \rightarrow \mathbb{R}$

Initialize value-function weights \mathbf{w} as appropriate (e.g., $\mathbf{w} = \mathbf{0}$)

Repeat forever:

 Generate an episode $S_0, A_0, R_1, S_1, A_1, \dots, R_T, S_T$ using π

 For $t = 0, 1, \dots, T - 1$:


$$\mathbf{w} \leftarrow \mathbf{w} + \alpha [G_t - \hat{v}(S_t, \mathbf{w})] \nabla \hat{v}(S_t, \mathbf{w})$$

\mathbf{w} are the tunable
parameters of the value
approximation function

- Guaranteed to converge to a local optimum because G_t is an unbiased estimate of $v_\pi(S_t)$

Example


$$f(S) = [2, 2, 1]$$



10

- $S = \{f_1(S), f_2(S), f_3(S)\}$
 - $f_{1,2}$ =distance to ghost 1,2, f_3 =distance to food
- $\hat{v}(s) = \sum_i \theta_i f_i(s)$
 - init: $\theta = [0, 0, 0]$
- $\theta = \theta + \alpha(G_t - \hat{v}(s; \theta)) \nabla \hat{v}(s; \theta)$
- $\theta = [0, 0, 0] + 0.1(10 - [0, 0, 0] \cdot [2, 2, 1])[2, 2, 1]$
 - $\theta = [2, 2, 1]$
- $\hat{v}(S') = f(S') \cdot \theta = [2, 4, 1] \cdot [2, 2, 1] = 13$

$$f(S') = [2, 4, 1]$$



Side note

- Should we care about on-policy value approximation?
 - Once the policy changes the approximated values become irrelevant
- Yes! This will be useful for Actor-Critic methods which will be discussed later



Learning approximation with bootstrapping

- Can we update the value approximation function at every step?
- Yes, define SGD as a function of the TD error
 - Tabular TD learning: $\hat{v}(s_t) = \hat{v}(s_t) + \alpha(r_t + \gamma\hat{v}(s_{t+1}) - \hat{v}(s_t))$
 - Approximation TD learning: $\theta = \theta + \alpha(r_t + \gamma\hat{v}(s_{t+1}; \theta) - \hat{v}(s_t; \theta))\nabla\hat{v}(s_t; \theta)$
- Known as Semi-gradient methods
- **NOT** guaranteed to converge to a local optimum because $\hat{v}(s_{t+1}; \theta)$ is a biased estimate of $v_\pi(s_{t+1})$
- Semi-gradient (bootstrapping) methods do not converge as robustly as (full) gradient methods

Semi-gradient methods

- They do converge reliably in important cases such as the linear approximation case
- They offer important advantages that make them often clearly preferred
- They typically enable significantly faster learning, as we have seen in Chapters 6 and 7
- They enable learning to be continual and online, without waiting for the end of an episode
- This enables them to be used on continuing problems and provides computational advantages

Semi-gradient TD(0)

Semi-gradient TD(0) for estimating $\hat{v} \approx v_\pi$

Input: the policy π to be evaluated

Input: a differentiable function $\hat{v} : \mathcal{S}^+ \times \mathbb{R}^d \rightarrow \mathbb{R}$ such that $\hat{v}(\text{terminal}, \cdot) = 0$

Initialize value-function weights \mathbf{w} arbitrarily (e.g., $\mathbf{w} = \mathbf{0}$)

Repeat (for each episode):

 Initialize S

 Repeat (for each step of episode):

 Choose $A \sim \pi(\cdot | S)$

 Take action A , observe R, S'

$\mathbf{w} \leftarrow \mathbf{w} + \alpha [R + \gamma \hat{v}(S', \mathbf{w}) - \hat{v}(S, \mathbf{w})] \nabla \hat{v}(S, \mathbf{w})$

$S \leftarrow S'$

 until S' is terminal

What's the difference
from the tabular case?

Semi-gradient TD(0)

Semi-gradient TD(0) for estimating $\hat{v} \approx v_\pi$

Input: the policy π to be evaluated

Input: a differentiable function $\hat{v} : \mathcal{S}^+ \times \mathbb{R}^d \rightarrow \mathbb{R}$ such that $\hat{v}(\text{terminal}, \cdot) = 0$

Initialize value-function weights \mathbf{w} arbitrarily (e.g., $\mathbf{w} = \mathbf{0}$)

Repeat (for each episode):

Initialize S

Repeat (for each step of episode):

Choose $A \sim \pi(\cdot|S)$

Take action A , observe R, S'

$\mathbf{w} \leftarrow \mathbf{w} + \alpha [R + \gamma \hat{v}(S', \mathbf{w}) - \hat{v}(S, \mathbf{w})] \nabla \hat{v}(S, \mathbf{w})$

$S \leftarrow S'$


until S' is terminal

What's the difference
from the tabular case?

$$\hat{v}(S) = \hat{v}(S) + \alpha [R + \gamma \hat{v}(S') - \hat{v}(S)]$$

Example


$$f(S) = [2,3,1]$$



$R = +10$




$$f(S') = [1,2,1]$$



- $S = \{f_1(S), f_2(S), f_3(S)\}$
 - $f_{1,2}$ =distance to ghost 1,2, f_3 =distance to food
- $\hat{v}(s) = \sum_i \theta_i f_i(s)$
 - init: $\theta = [0,0,0]$
- $\theta = \theta + \alpha(R + \gamma \hat{v}(S'; \theta) - \hat{v}(S; \theta)) \nabla \hat{v}(S; \theta)$
- $\theta = [0,0,0] + 0.1(10 + [1,2,1] \cdot [0,0,0] - [2,3,1] \cdot [0,0,0])[2,3,1]$
 - $\theta = [2,3,1]$
- $\hat{v}(U) = f(U) \cdot \theta = [2,4,1] \cdot [2,3,1] = 17$

$$f(U) = [2,4,1]$$



n -step return

n -step semi-gradient TD for estimating $\hat{v} \approx v_\pi$

Input: the policy π to be evaluated

Input: a differentiable function $\hat{v} : \mathcal{S}^+ \times \mathbb{R}^d \rightarrow \mathbb{R}$ such that $\hat{v}(\text{terminal}, \cdot) = 0$

Parameters: step size $\alpha \in (0, 1]$, a positive integer n

All store and access operations (S_t and R_t) can take their index mod n

Initialize value-function weights \mathbf{w} arbitrarily (e.g., $\mathbf{w} = \mathbf{0}$)

Repeat (for each episode):

 Initialize and store $S_0 \neq \text{terminal}$

$T \leftarrow \infty$

 For $t = 0, 1, 2, \dots$:

 If $t < T$, then:

 Take an action according to $\pi(\cdot | S_t)$

 Observe and store the next reward as R_{t+1} and the next state as S_{t+1}

 If S_{t+1} is terminal, then $T \leftarrow t + 1$

$\tau \leftarrow t - n + 1$ (τ is the time whose state's estimate is being updated)

 If $\tau \geq 0$:

$G \leftarrow \sum_{i=\tau+1}^{\min(\tau+n, T)} \gamma^{i-\tau-1} R_i$

 If $\tau + n < T$, then: $G \leftarrow G + \gamma^n \hat{v}(S_{\tau+n}, \mathbf{w})$ ($G_{\tau:\tau+n}$)

$\mathbf{w} \leftarrow \mathbf{w} + \alpha [G - \hat{v}(S_\tau, \mathbf{w})] \nabla \hat{v}(S_\tau, \mathbf{w})$

 Until $\tau = T - 1$

- Again, only a simple modification over the tabular setting

n -step return

n -step semi-gradient TD for estimating $\hat{v} \approx v_\pi$

Input: the policy π to be evaluated

Input: a differentiable function $\hat{v} : \mathcal{S}^+ \times \mathbb{R}^d \rightarrow \mathbb{R}$ such that $\hat{v}(\text{terminal}, \cdot) = 0$

Parameters: step size $\alpha \in (0, 1]$, a positive integer n

All store and access operations (S_t and R_t) can take their index mod n

Initialize value-function weights \mathbf{w} arbitrarily (e.g., $\mathbf{w} = \mathbf{0}$)

Repeat (for each episode):

Initialize and store $S_0 \neq \text{terminal}$

$T \leftarrow \infty$

For $t = 0, 1, 2, \dots$:

 If $t < T$, then:

 Take an action according to $\pi(\cdot | S_t)$

 Observe and store the next reward as R_{t+1} and the next state as S_{t+1}

 If S_{t+1} is terminal, then $T \leftarrow t + 1$

$\tau \leftarrow t - n + 1$ (τ is the time whose state's estimate is being updated)

 If $\tau \geq 0$:

$G \leftarrow \sum_{i=\tau+1}^{\min(\tau+n, T)} \gamma^{i-\tau-1} R_i$

 If $\tau + n < T$, then: $G \leftarrow G + \gamma^n \hat{v}(S_{\tau+n}, \mathbf{w})$ ($G_{\tau:\tau+n}$)


$\mathbf{w} \leftarrow \mathbf{w} + \alpha [G - \hat{v}(S_\tau, \mathbf{w})] \nabla \hat{v}(S_\tau, \mathbf{w})$

Until $\tau = T - 1$

- Again, only a simple modification over the tabular setting
- Weight update instead of tabular entry update

Another optimization approach

- Approach1: Gradient decent
 - Update θ in iterations
 - Find a fixed point
 - If: $\theta_{t+1} = \theta_t$
 - Then: converged to a local optimum
- Approach2: Compute the fixed point directly
 - Solve: $\theta_{t+1} = \theta_t$



Compute fixed point over θ

- Assume a linear function approximator $\hat{v}(f(s); \theta) = f(s) \cdot \theta$
- TD update: $\theta_{t+1} = \theta_t + \alpha [R_{t+1} + \gamma \hat{v}(S_{t+1}) - \hat{v}(S_t)] \nabla \hat{v}(S_t)$
- $= \theta_t + \alpha [R_{t+1} + (\gamma f(S_{t+1}) - f(S_t)) \theta_t] f(S_t)$
- $= \theta_t + \alpha (\mathbf{b} - \mathbf{A} \theta_t)$
- Where: $\mathbf{b} = \mathbb{E}[R_{t+1} f(S_t)]$, $\mathbf{A} = \mathbb{E} [f(S_t) (f(S_t) - \gamma f(S_{t+1}))^\top]$
- Fixed point at: $\mathbb{E}[\theta_{t+1}] = \mathbb{E}[\theta_t]$
 - $\mathbf{b} - \mathbf{A} \theta_t = 0$
 - $\mathbf{b} = \mathbf{A} \theta_t$
 - $\theta_t = \mathbf{A}^{-1} \mathbf{b}$

For a linear approximator

TD-error = 0

Least squares TD

- Approximate A and b online, solve $\theta = A^{-1}b$
- Where: $b = \mathbb{E}[R_{t+1}f(S_t)]$, $A = \mathbb{E}\left[f(S_t)(f(S_t) - \gamma f(S_{t+1}))^\top\right]$
- $\hat{A} = \sum_{k=0}^{t-1} f(S_k)(f(S_k) - \gamma f(S_{k+1}))^\top$
- But \hat{A} is not guaranteed to be invertible (it might have '0' on diagonal)
- So add a small constant (ε) to the diagonal
 - $\hat{A} = \sum_t f(S_t)(f(S_t) - \gamma f(S_{t+1}))^\top + \varepsilon I$
- $\hat{b} = \sum_t R_{t+1}f(S_t)$

Least squares TD

LSTD for estimating $\hat{v} \approx v_\pi$ ($O(d^2)$ version)

Input: feature representation $\mathbf{x}(s) \in \mathbb{R}^d$, for all $s \in \mathcal{S}$, $\mathbf{x}(\text{terminal}) \doteq \mathbf{0}$

$$\widehat{\mathbf{A}}^{-1} \leftarrow \varepsilon^{-1} \mathbf{I}$$

$$\widehat{\mathbf{b}} \leftarrow \mathbf{0}$$

Store the inverse of A
instead of A

An $d \times d$ matrix

An d -dimensional vector

Repeat (for each episode):

Initialize S ; obtain corresponding \mathbf{x}

Repeat (for each step of episode):

Choose $A \sim \pi(\cdot|S)$

Take action A , observe R, S' ; obtain corresponding \mathbf{x}'

$$\mathbf{v} \leftarrow \widehat{\mathbf{A}}^{-1} (\mathbf{x} - \gamma \mathbf{x}')$$

$$\widehat{\mathbf{A}}^{-1} \leftarrow \widehat{\mathbf{A}}^{-1} - (\widehat{\mathbf{A}}^{-1} \mathbf{x}) \mathbf{v}^\top / (1 + \mathbf{v}^\top \mathbf{x})$$

$$\widehat{\mathbf{b}} \leftarrow \widehat{\mathbf{b}} + R \mathbf{x}$$

$$\boldsymbol{\theta} \leftarrow \widehat{\mathbf{A}}^{-1} \widehat{\mathbf{b}}$$

$$S \leftarrow S'; \mathbf{x} \leftarrow \mathbf{x}'$$

until S' is terminal

Least squares TD

LSTD for estimating $\hat{v} \approx v_\pi$ ($O(d^2)$ version)

Input: feature representation $\mathbf{x}(s) \in \mathbb{R}^d$, for all $s \in \mathcal{S}$, $\mathbf{x}(\text{terminal}) \doteq \mathbf{0}$

$$\widehat{\mathbf{A}}^{-1} \leftarrow \varepsilon^{-1} \mathbf{I}$$

An $d \times d$ matrix

$$\widehat{\mathbf{b}} \leftarrow \mathbf{0}$$

An d -dimensional vector

Repeat (for each episode):

Initialize S ; obtain corresponding \mathbf{x}

Repeat (for each step of episode):

Choose $A \sim \pi(\cdot|S)$

Take action A , observe R, S' ; obtain corresponding \mathbf{x}'

$$\mathbf{v} \leftarrow \widehat{\mathbf{A}}^{-1} (\mathbf{x} - \gamma \mathbf{x}')$$

$$\widehat{\mathbf{A}}^{-1} \leftarrow \widehat{\mathbf{A}}^{-1} - (\widehat{\mathbf{A}}^{-1} \mathbf{x}) \mathbf{v}^\top / (1 + \mathbf{v}^\top \mathbf{x})$$

$$\widehat{\mathbf{b}} \leftarrow \widehat{\mathbf{b}} + R \mathbf{x}$$

$$\boldsymbol{\theta} \leftarrow \widehat{\mathbf{A}}^{-1} \widehat{\mathbf{b}}$$

$$S \leftarrow S'; \mathbf{x} \leftarrow \mathbf{x}'$$


until S' is terminal

Incremental updates
(no need to store all
previous transitions)



Least squares TD

- Directly computing the TD fixed point
- Most data efficient form of linear TD(0), but it is also more expensive computationally
 - Semi-gradient linear TD(0) requires memory and per-step computation that is only $O(d)$ where d is the number of state features
- In the incremental update version, \hat{A} is an outer product (a column vector times a row vector) and thus requires a matrix update
- The update computational complexity is $O(d^2)$, and the memory complexity is $O(d^2)$

Feature selection




- Assume a linear function approximator $\hat{v}(f(s); \theta) = f(s) \cdot \theta$
- What relevant features should represent states?



Features are domain depended
requiring expert knowledge

Automatic features extraction

- Consider a game state as a pixel matrix
- Raw data of type: $pixel(7,3) = [0,0,0]$ (black)
- Desired features = {ball location, ball speed, ball direction, pan location...}
- How can we translate pixels to the relevant features?




Automatic features extraction for linear approximator

- **Polynomials:** $f_i(s) = \prod_{j=1}^k x_j^{c_{i,j}}$
 - where each $c_{i,j}$ is an integer in the set $\{0, 1, \dots, n\}$ for an integer $n \geq 0$
 - These features makeup the order- n polynomial basis for dimension k , which contains $(n + 1)^k$ different features
- **Fourier Basis:** $f_i(s) = \cos(\pi X^\top c^i)$
 - Where $c^i = (c_1^i, \dots, c_k^i)^\top$, with $c_j^i \in \{0, \dots, n\}$ for $j = \{1, \dots, k\}$ and $i = \{0, \dots, (n + 1)^k\}$
 - This defines a feature for each of the $(n + 1)^k$ possible integer vectors c^i
 - The inner product $X^\top c^i$ has the effect of assigning an integer in $\{0, \dots, n\}$ to each dimension of X
 - This integer determines the feature's frequency along that dimension
 - The features can be shifted and scaled to suit the bounded state space of a particular application

Automatic features extraction for linear approximator

- Other approaches include: Coarse Coding, Tile Coding, Radial Basis Functions (See chapter 9.5 in textbook)
- Each of these approaches defines a set of features, some useful yet most are not
 - E.g., is there a polynomial/Fourier function that translates pixels to pan location?
 - Probably but it's a needle in a (combinatorial) haystack
- Can we do better (generically)
 - Yes, using deep neural networks...



What did we learn?

- Reinforcement learning must generalize on observed experience if it is to be applicable to real world domains
- We can use parameterized function approximation to represent our knowledge about the domain state/action values
- Use stochastic gradient descent to update the tunable parameters such that the observed (TD, rollout) error is reduced
- When using a linear approximator, the Least squares TD method provides the most sample efficient approximation

What next?

- **Lecture:** Deep Neural Networks as function approximators
- **Assignments:**
 - Value Iteration, by September-23, EOD
 - Asynchronous Value Iteration, by September-23, EOD
 - Policy Iteration, by September-23, EOD
 - Monte-Carlo Control by September-30, EOD
 - Monte-Carlo Control with Importance Sampling by September-30, EOD
 - Tabular Q-Learning, by Oct. 7, EOD
 - SARSA, by Oct. 7, EOD
 - Q-Learning with Approximation, by Oct. 14, EOD
- **Quiz (on Canvas):**
 - n-step Bootstrapping, by Sep. 23, EOD
 - Value approximation, by Oct. ?, EOD
- **Project:**
 - Project proposal, by Sep. 30 EOD