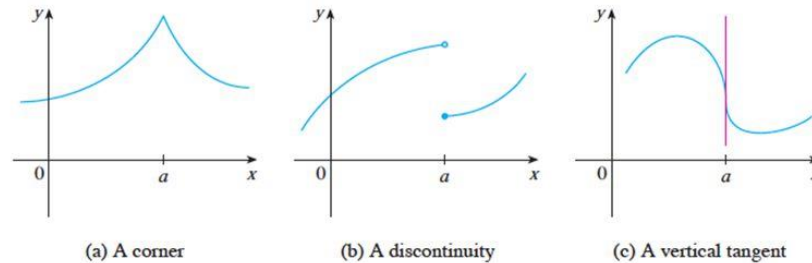


CSCE-642 Reinforcement Learning

Derivative Free Methods

Non-differentiable function

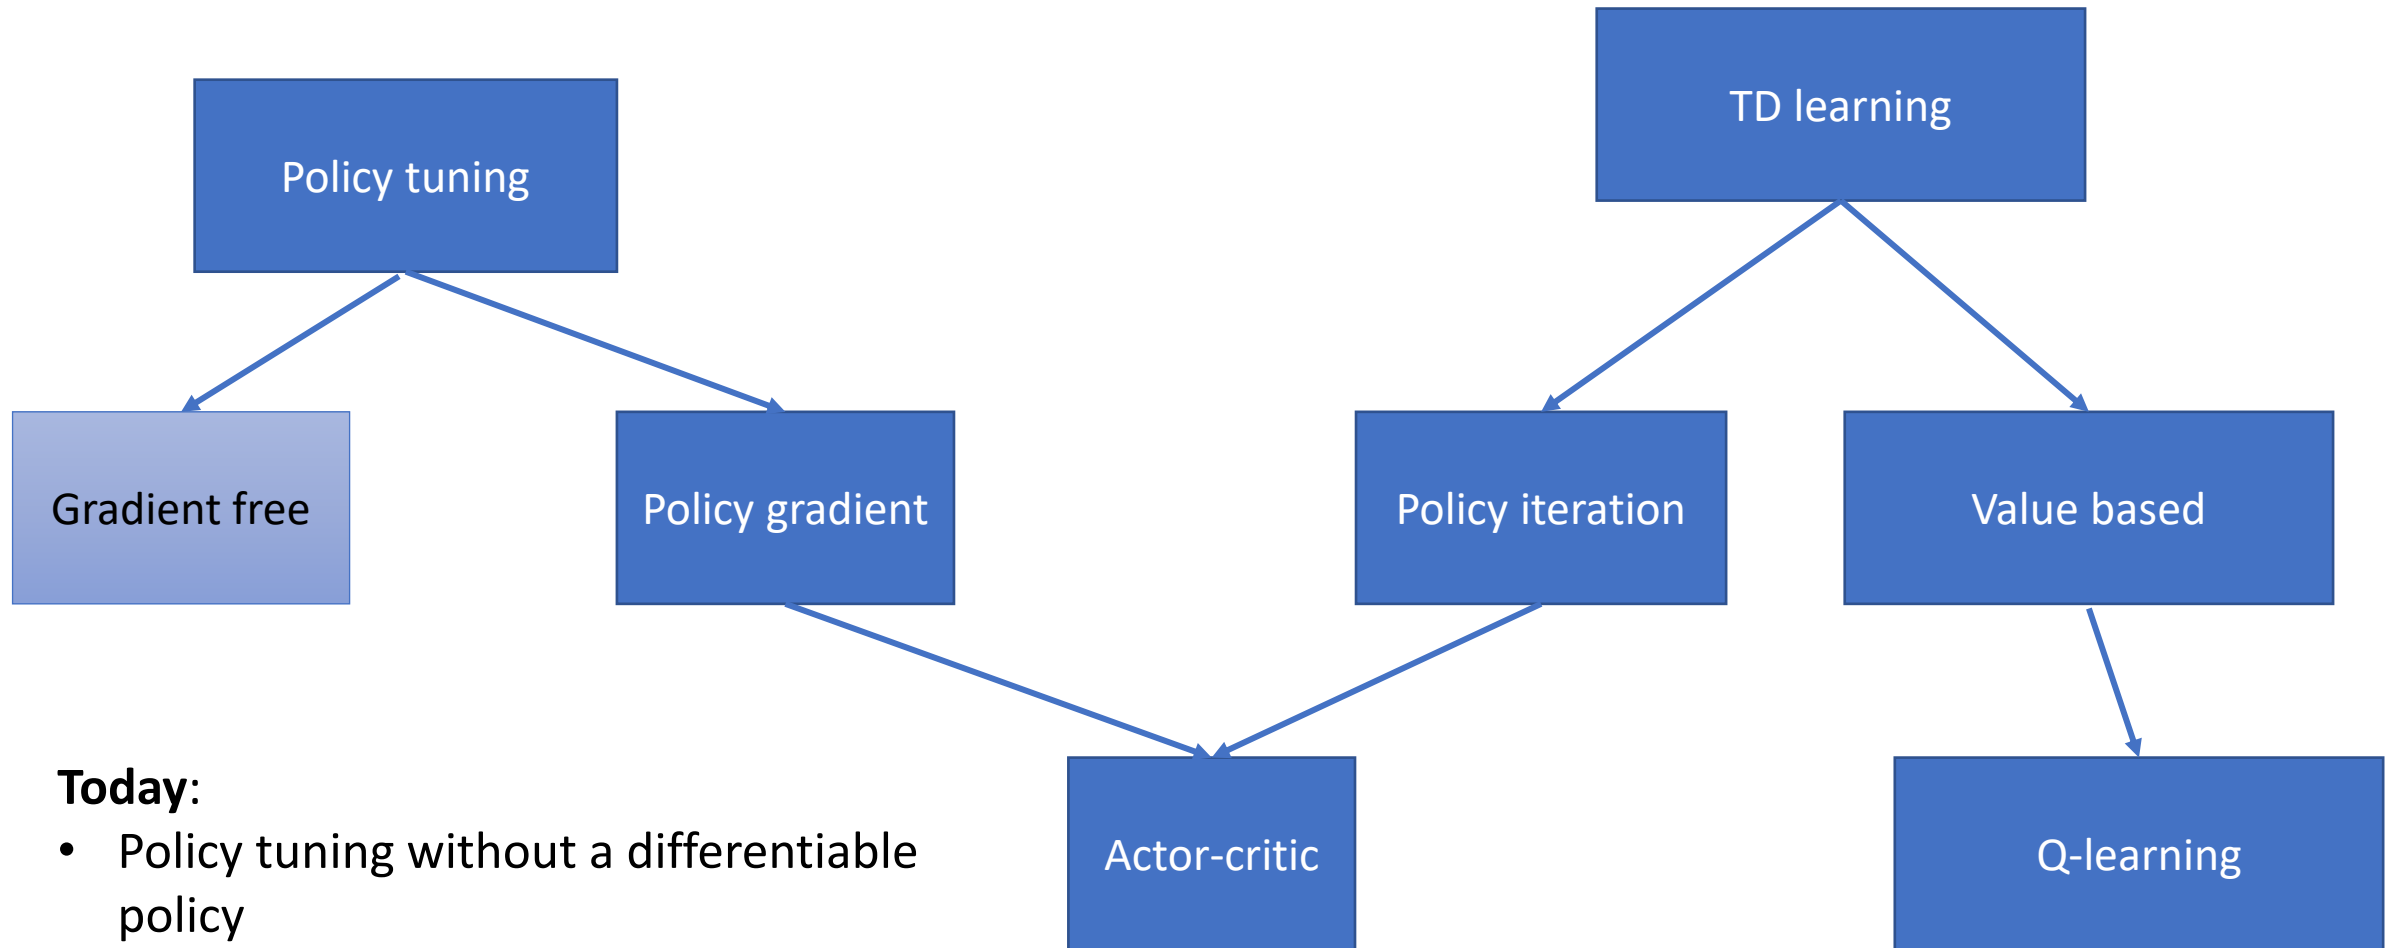


Instructor: Guni Sharon

Final report

- Please limit to 6 pages
- Be concise and on-point
- Follow a clear narrative

Solver classes



Today:

- Policy tuning without a differentiable policy
- Also known as gradient free methods

Online parameter tuning

1. Gradient approximation -

Finite Difference Policy Gradient Decent

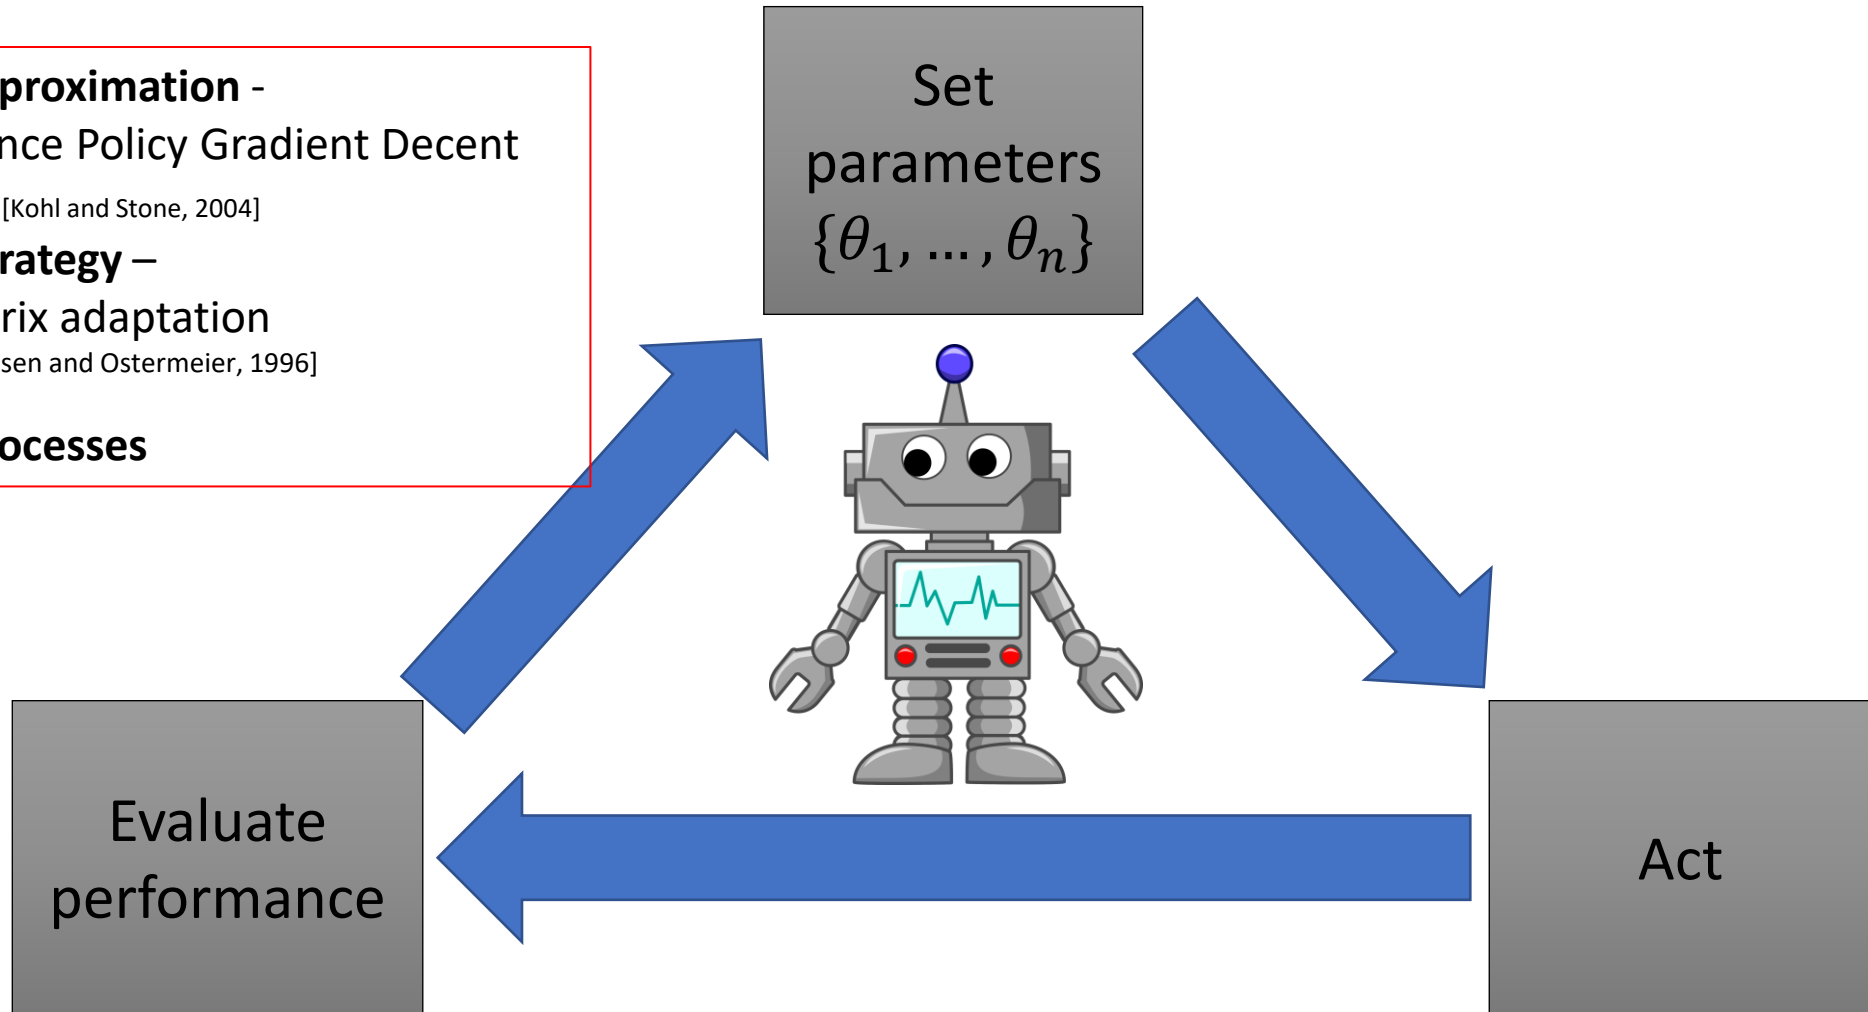
[Kohl and Stone, 2004]

2. Evolution strategy –

covariance matrix adaptation

[Hansen and Ostermeier, 1996]

3. Gaussian processes



Online parameter tuning

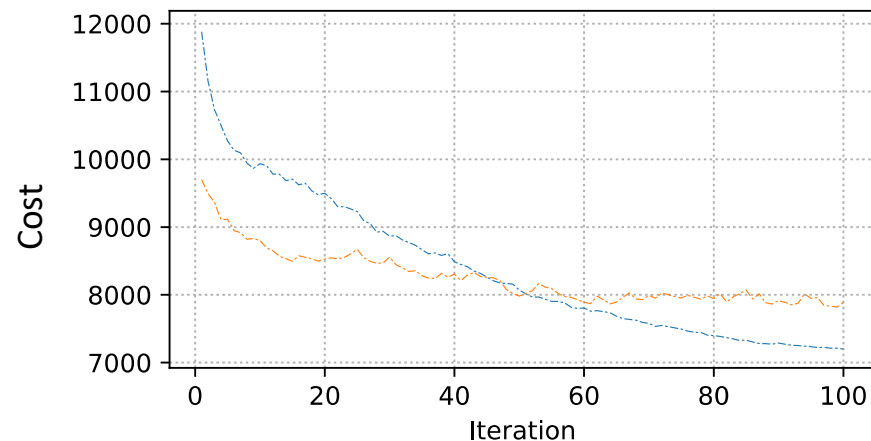
- Minimize $f: \mathbb{R}^n \rightarrow \mathbb{R}$
 - Where is the optimum of $f(X)$?
 - If convex then at $\nabla_X f = 0$ (Lagrange, KKT), else, gradient decent (ADAM, Newton)
- What if f is unknown and we can only sample $f(X)$ for a given X ?
 - Approximate f with supervised learning
- Observations are not provided -> must learn from interactions
 - Reinforcement learning!
- f is unknown/not differentiable
 - Blackbox (derivative free) optimization!

Examples

- **Comparison between popular genetic algorithm (GA)-based tool and covariance matrix adaptation – evolutionary strategy (CMA-ES) for optimizing indoor daylight**
Manal Anis, Sumedh Pendurkar, Yun Kyu Yi, and Guni Sharon
In *Proceedings of Building Simulation 2023: 18th Conference of IBPSA*, 2023
- **Bilevel Entropy based Mechanism Design for Balancing Meta in Video Games.**
Sumedh Pendurkar, Chris Chow, Luo Jie, and Guni Sharon
In *Proceedings of the 22th International Conference on Autonomous Agents and MultiAgent Systems (AAMAS)*, 2023

Blackbox optimization

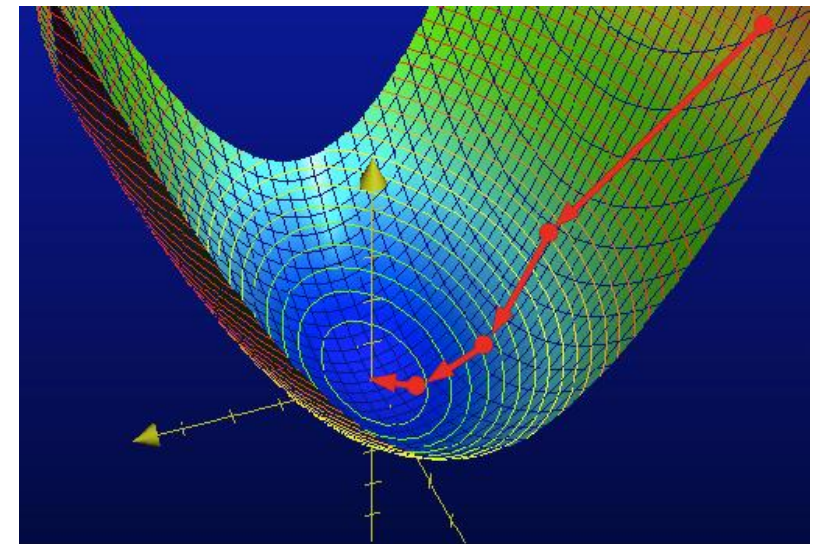
- Challenges:
 - Compute the optimum of an unknown function via samples
 - **Sample efficiency**: minimize the number of samples along the training process
 - **minimize** cumulative regret (for online optimization)



1. Gradient Descent

- Minimize $f(x_1, x_2, \dots, x_n)$
 1. Set initial parameter vector $X^0 = [x_1^0, \dots, x_n^0]^\top$
 2. While improving
 - $X^{k+1} = X^k - \alpha \nabla f^k$?

How can we compute
the gradient if f is
unknown?



Empirical gradient

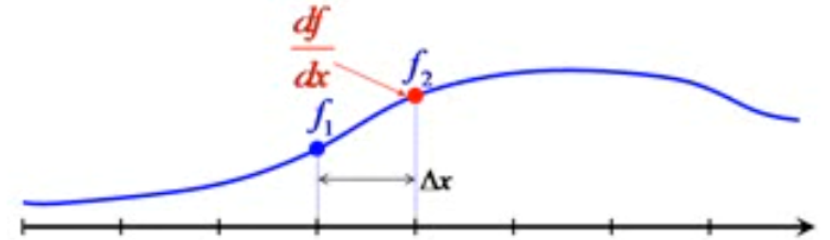
- approximate $\nabla f(x_1, x_2, \dots, x_n)$ at X
 - Through local evaluation
- Finite-difference methods (FDM)
 - By definition: $\frac{\partial f}{\partial x_i} = \lim_{x_i - x'_i \rightarrow 0} \frac{f(X) - f(X')}{x_i - x'_i}$
 - Not well-defined for non-differentiable f
 - However, we can avoid infinities by setting $x_i - x'_i > 0$, $\frac{\partial f}{\partial x_i} \cong \frac{f(X) - f(X')}{x_i - x'_i}$

Finite difference approach

Biased towards the rear

Backward
difference

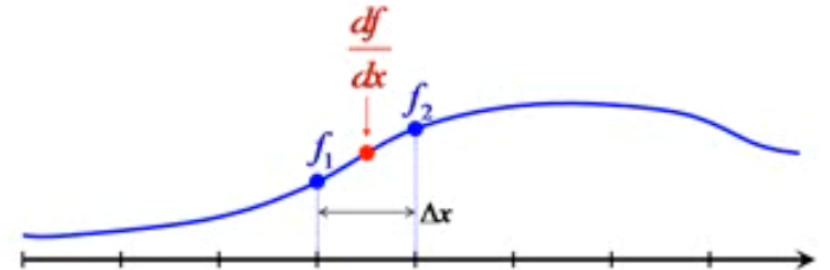
$$\frac{df_2}{dx} \approx \frac{f_2 - f_1}{\Delta x}$$



Unbiased
X2 samples

Central
difference

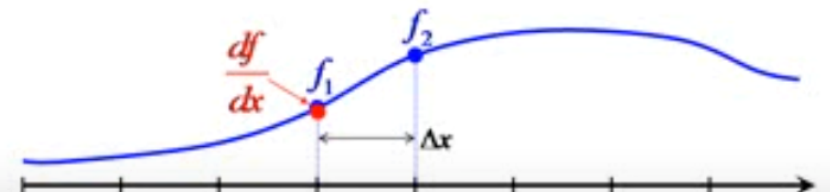
$$\frac{df_{1.5}}{dx} \approx \frac{f_2 - f_1}{\Delta x}$$



Biased towards the front

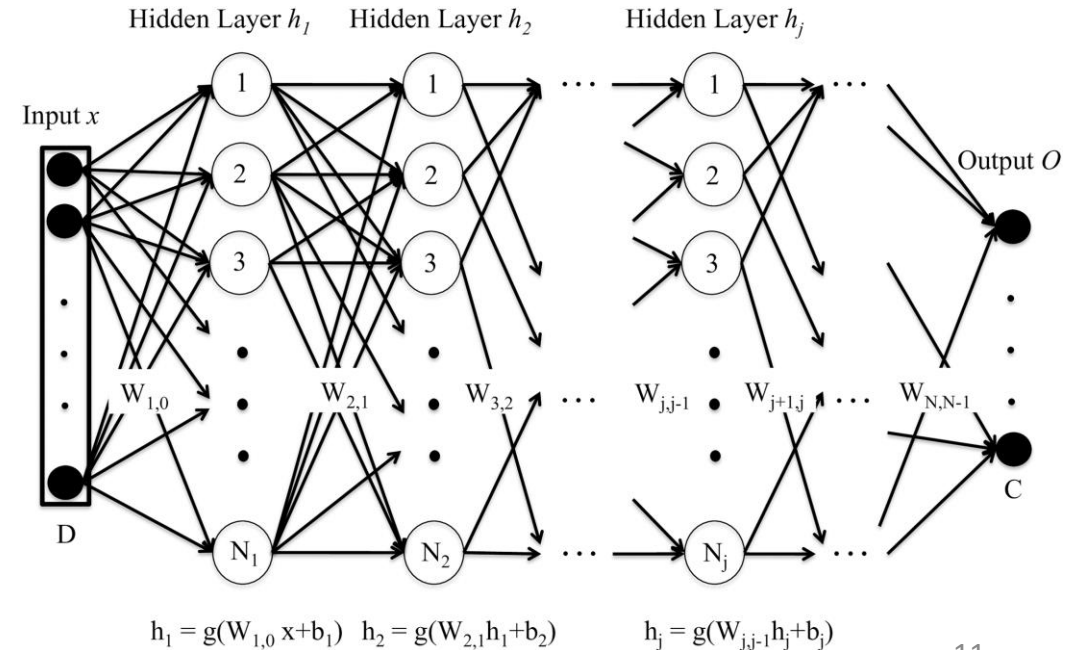
Forward
difference

$$\frac{df_1}{dx} \approx \frac{f_2 - f_1}{\Delta x}$$



High dimensionality

- Approximate the gradient of a DNN with 1M parameters
- Requires 2M observations!
 - Not practical
- A stochastic approach is required



Finite Difference Policy Gradient

(Kohl and Stone, 2004)

- Evaluate ∇f^k through M observations

- For $m = [1, \dots, M]$

- $\tilde{X}^{k,m} = \langle x_1^k + \delta_1^m, \dots, x_n^k + \delta_n^m \rangle^T$

- $c_{+\varepsilon n}^k = \text{Mean}_{x \in \tilde{X}_{+\varepsilon n}^k} (f(x))$

Similar definition for $c_{-\varepsilon n}^k$ and c_{0n}^k

$$\delta_n^m = \text{Rand}(-\varepsilon, 0, \varepsilon)$$

$\tilde{X}_{+\varepsilon n}^k$ = all parameter vectors where x_n^k was increased by ε

$$\bullet \frac{\partial f^k}{\partial x_n^k} = \begin{cases} 0 & , c_{+\varepsilon n}^k < c_{0n}^k \text{ \& } c_{-\varepsilon n}^k < c_{0n}^k \\ c_{+\varepsilon n}^k - c_{-\varepsilon n}^k , & \text{else} \end{cases}$$

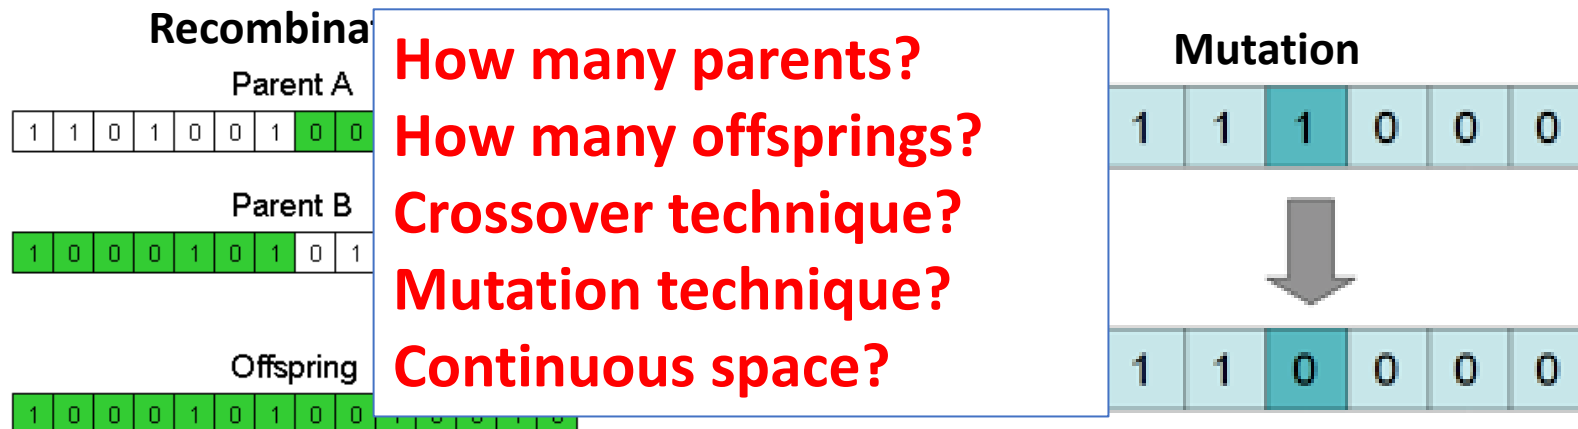
Example

- Minimize $f(x_1, x_2)$, $M = 4$, $\varepsilon = 1$, $\alpha = 1$
 - Initial parameter vector $X^0 = [0,0]^\top$
 - $\tilde{X}^{0,0} = [1,1]^\top, f(\tilde{X}^{0,0}) = 5$
 - $\tilde{X}^{0,1} = [-1,0]^\top, f(\tilde{X}^{0,1}) = 3$
 - $\tilde{X}^{0,2} = [0,0]^\top, f(\tilde{X}^{0,2}) = 6$
 - $\tilde{X}^{0,3} = [1,-1]^\top, f(\tilde{X}^{0,3}) = 4$
- $c_{+\varepsilon 1}^0, c_{-\varepsilon 1}^0, c_{01}^0, c_{+\varepsilon 2}^0, c_{-\varepsilon 2}^0, c_{02}^0 =$
- $\nabla f^0 =$
- $X^1 =$

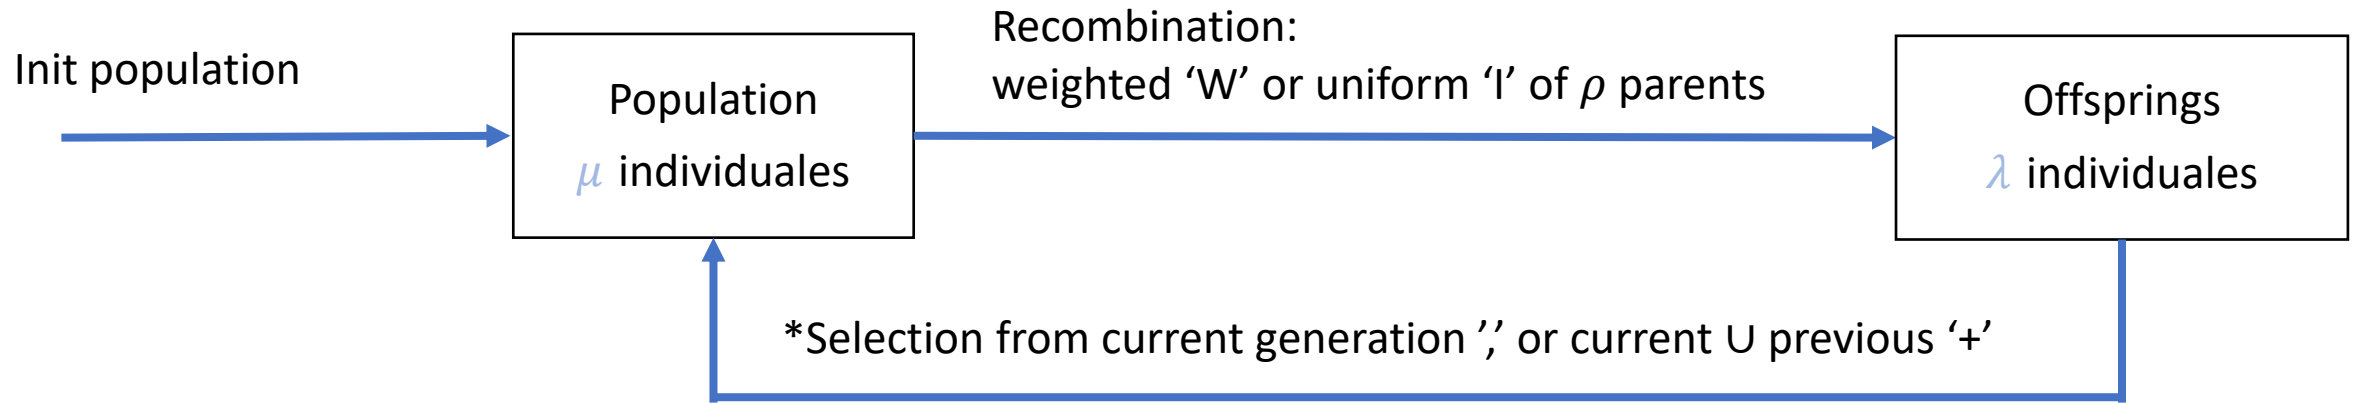
Genetic Algorithm

1. Generate M parameter vectors $\{\tilde{X}^{0,0}, \tilde{X}^{0,1}, \dots, \tilde{X}^{0,2}\}$
2. **Selection:** Survival of the fittest
 - $G \leftarrow$ set of parameter vectors with best fitness: $f(\tilde{X}^{k,m})$
3. Generate next generation through **recombination** and **mutation** on G
4. Goto 2. until termination criterion fulfilled

Might apply to the unification of current and former generations



$(\mu/\rho_{\{I,W\}}^+ , \lambda)$ -Evolutionary Strategy



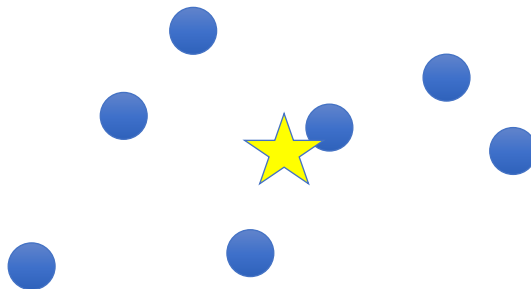
If comma selection, then $\lambda > \mu$

*While the ',' selection is recommended for unbounded search spaces (Schwefel, 1987), the '+' selection should be used in discrete finite size search spaces, e.g., in combinatorial optimization problems (Herdy, 1990; Beyer, 1992).

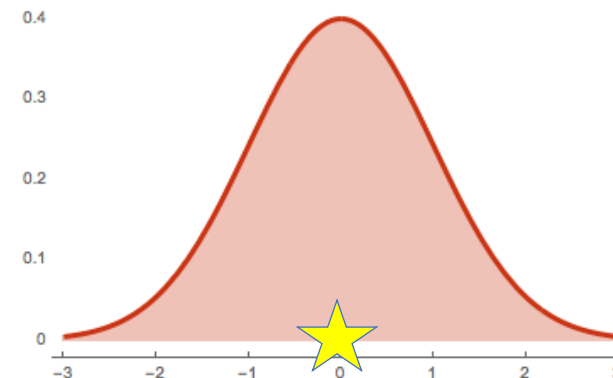
Continuous space

- **Recombination:** x_i for offspring = (weighted) mean m_i over the group ρ from the fittest individuals
- **Mutation:** for every offspring, sample each parameter $x_i \sim \mathcal{N}(m_i, Var_i)$

Recombination



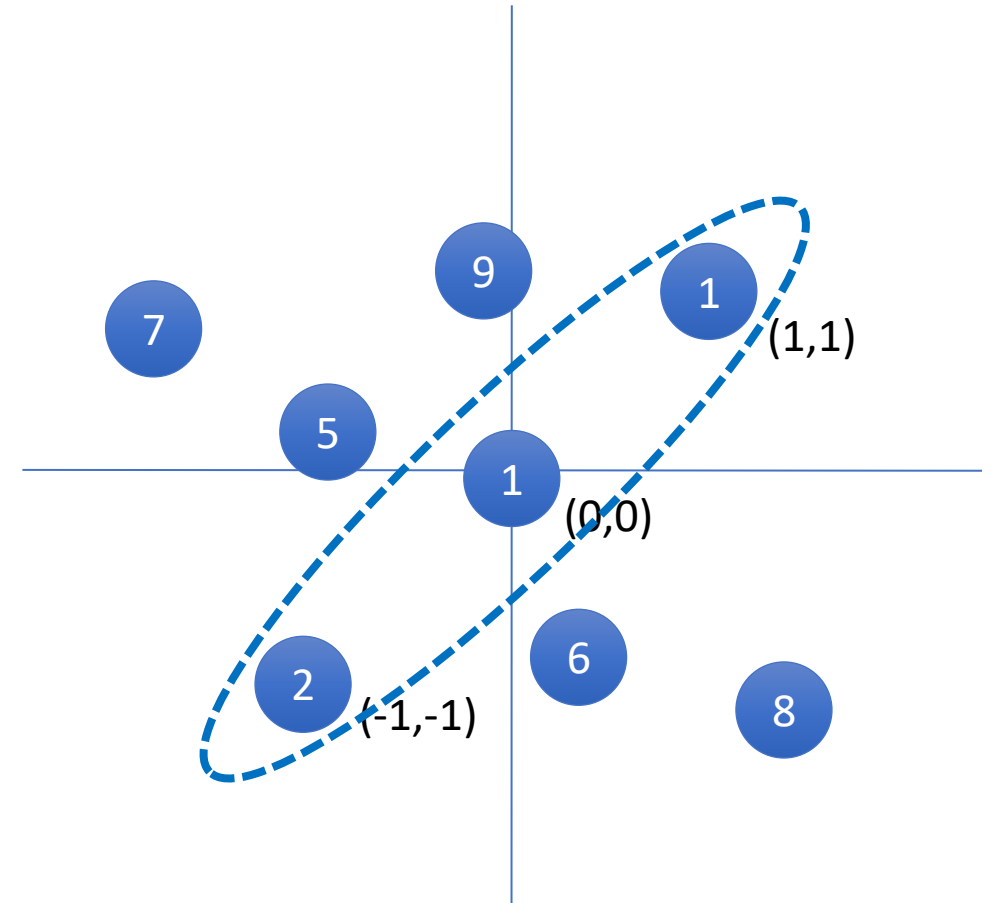
Mutation



Recombination

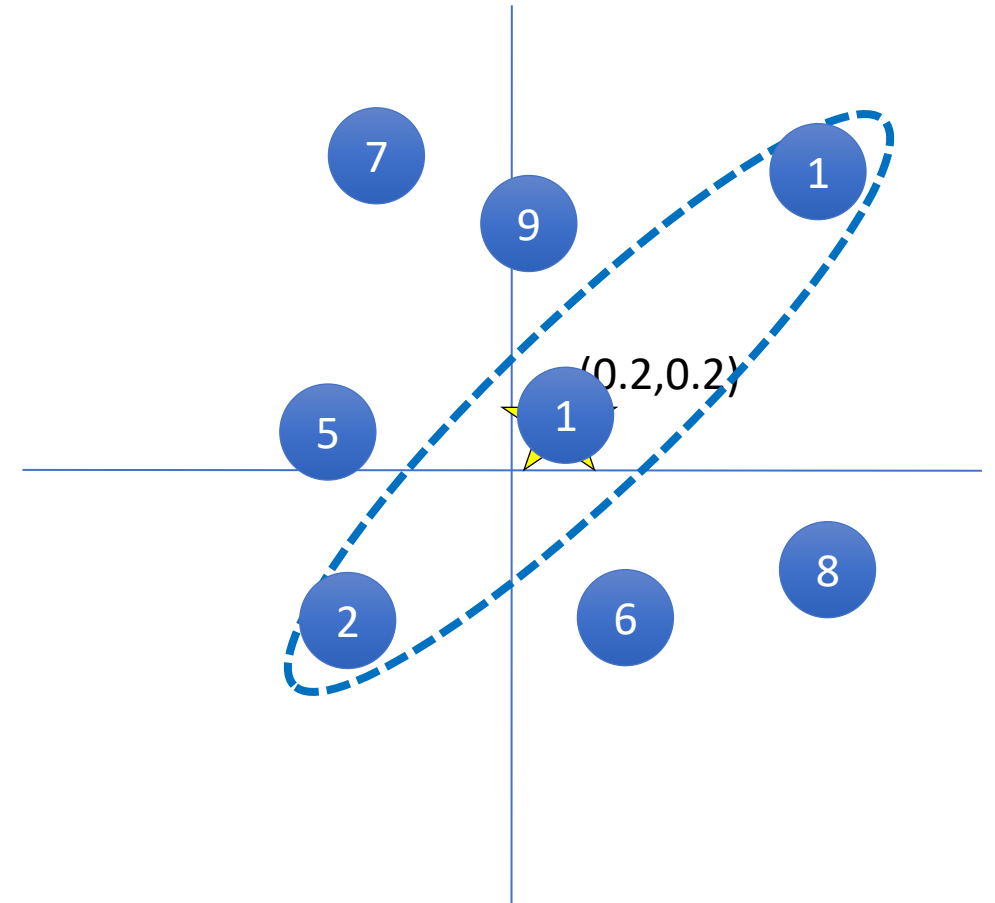
- Assume $(3/3_W, 8)$ -ES
- Select 3 fittest parents
- Recombine (weighted) groups of size 3
- Generate 8 offsprings
- Update:
 - $m_1 =$
 - $m_2 =$
 - $Var_1 =$
 - $Var_2 =$

$$\text{sample var} = \frac{\sum (x_i - \bar{x})^2}{n - 1}$$



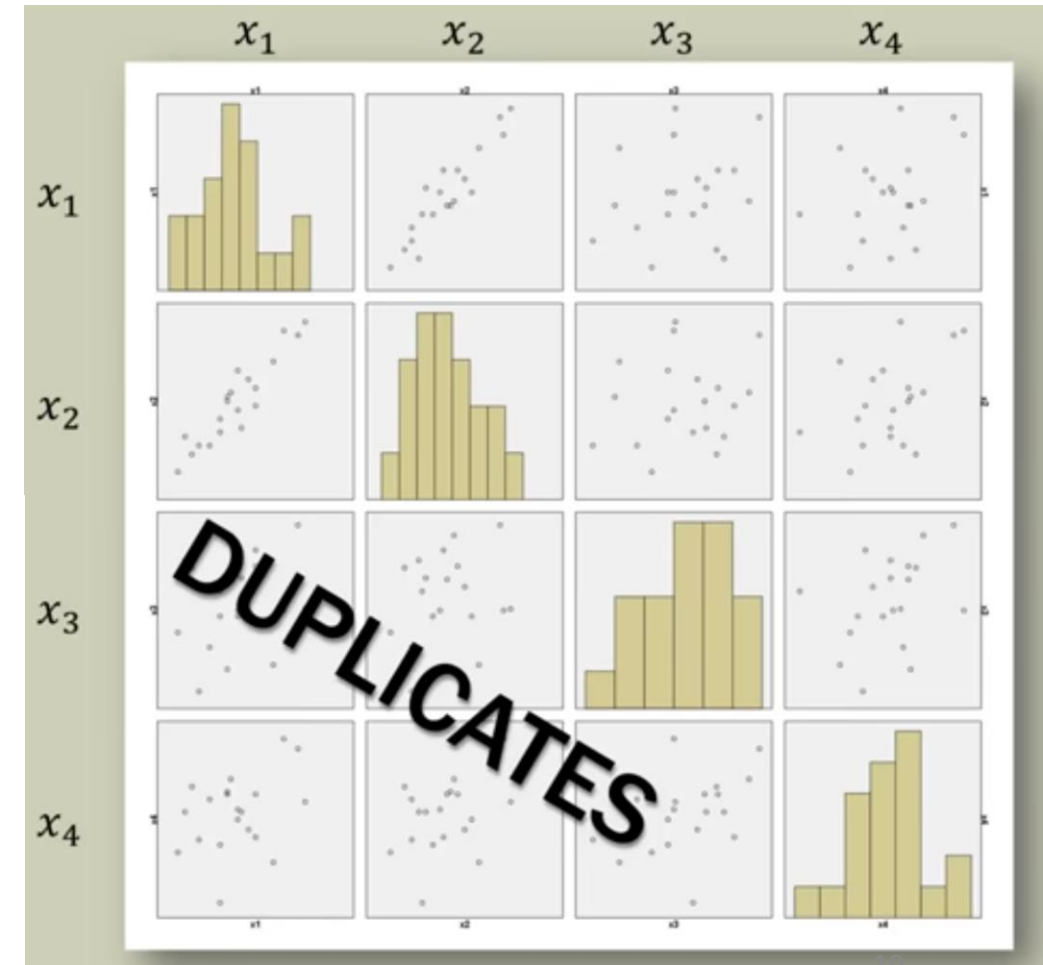
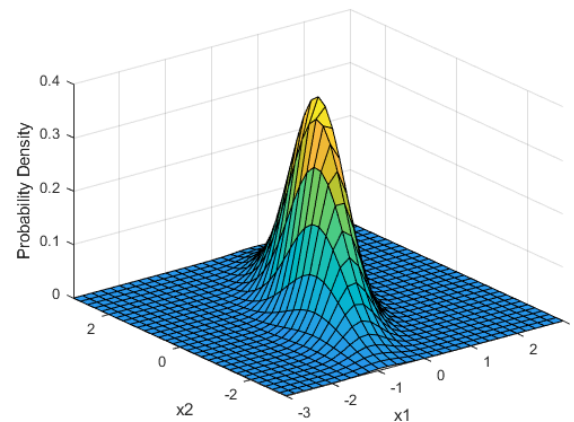
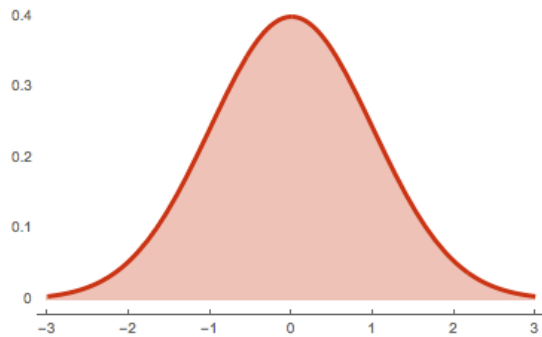
Mutation

- generate 8 offsprings
 - $x_1 \sim \mathcal{N}(0.2, 0.7)$
 - $x_2 \sim \mathcal{N}(0.2, 0.7)$
- Ooops! What went wrong?
 - We failed to enforce the variables' co-dependencies
 - Instead of Var_i consider the covariance matrix (C)



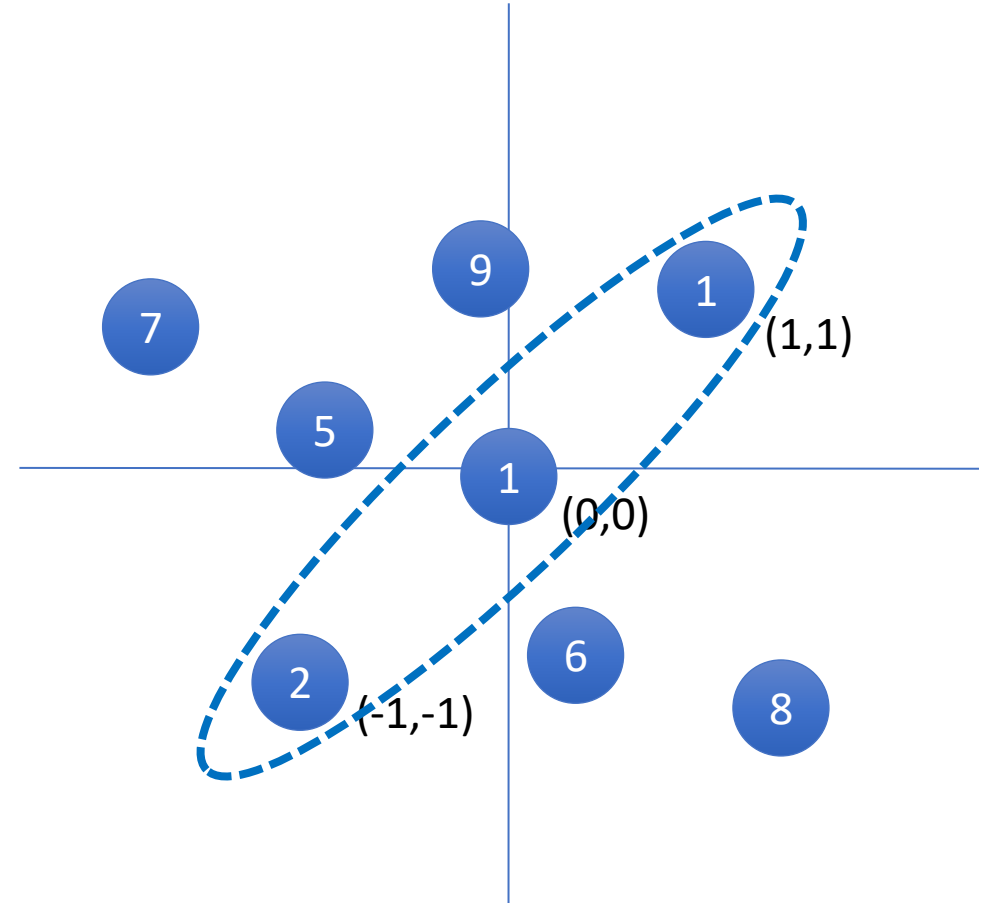
Covariance matrix

- Capture dependencies between the dimensions
- Enforce such dependencies when spawning offsprings
- $X \sim \mathcal{N}(\mu, C)$



Lets try again...

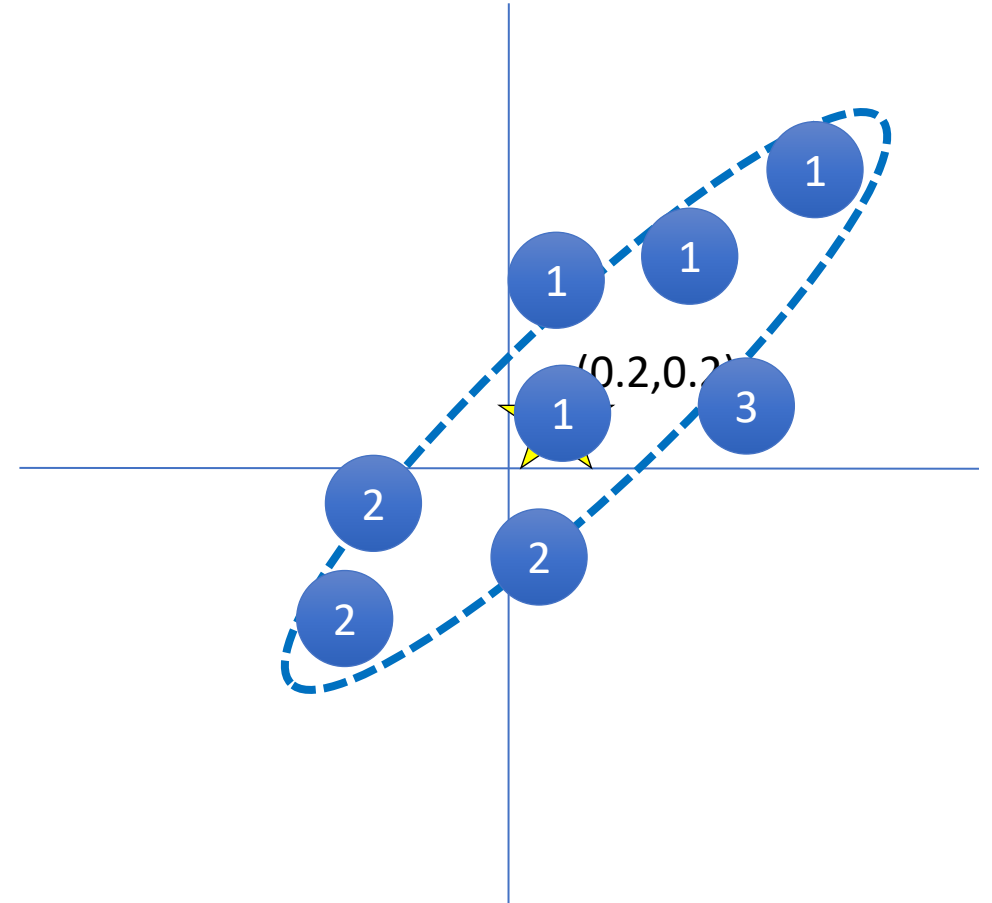
- Assume $(3/3_W, 8)$ CM-ES
- Select 3 fittest parents
- Update mean and covariance
 - $m =$
 - $C =$



$$*C_{ij} = \mathbb{E}[(x_i - m_i)(x_j - m_j)] = \frac{1}{n-1} \sum_{k=1}^M (x_i^k - m_i)(x_j^k - m_j)$$

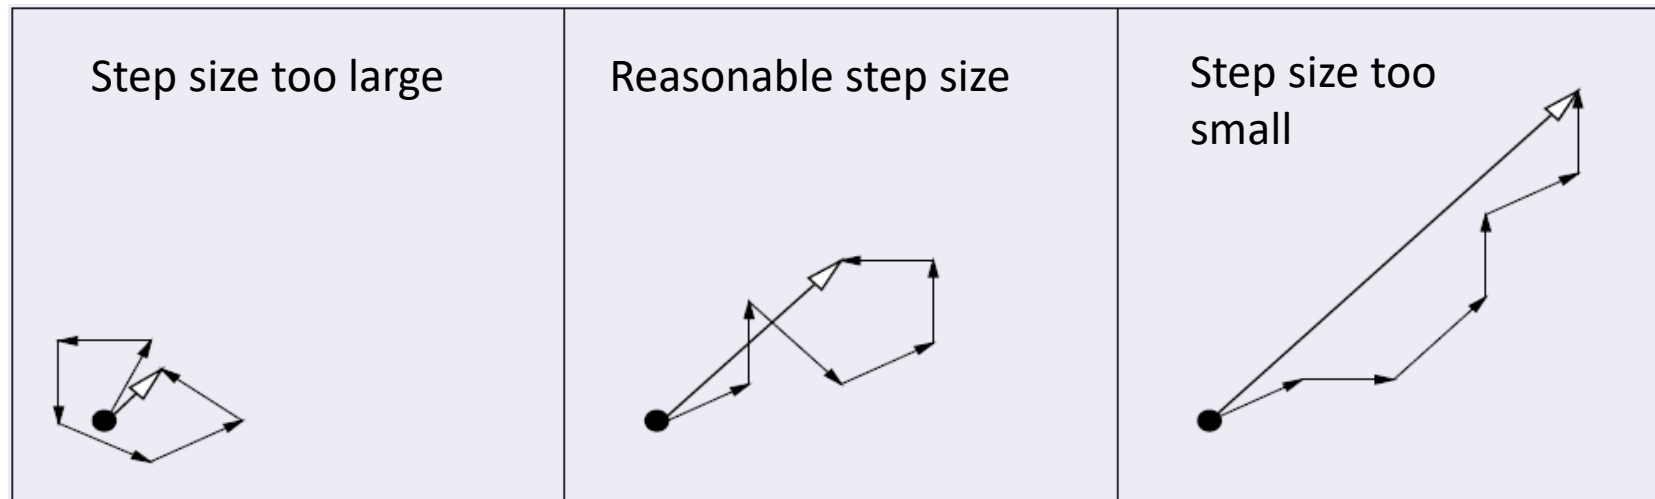
Mutation

- Create 8 offsprings
 - $X \sim \mathcal{N} \left(\begin{bmatrix} 0.2 \\ 0.2 \end{bmatrix}, \begin{bmatrix} 0.7 & 0.7 \\ 0.7 & 0.7 \end{bmatrix} \right)$
- Much better!



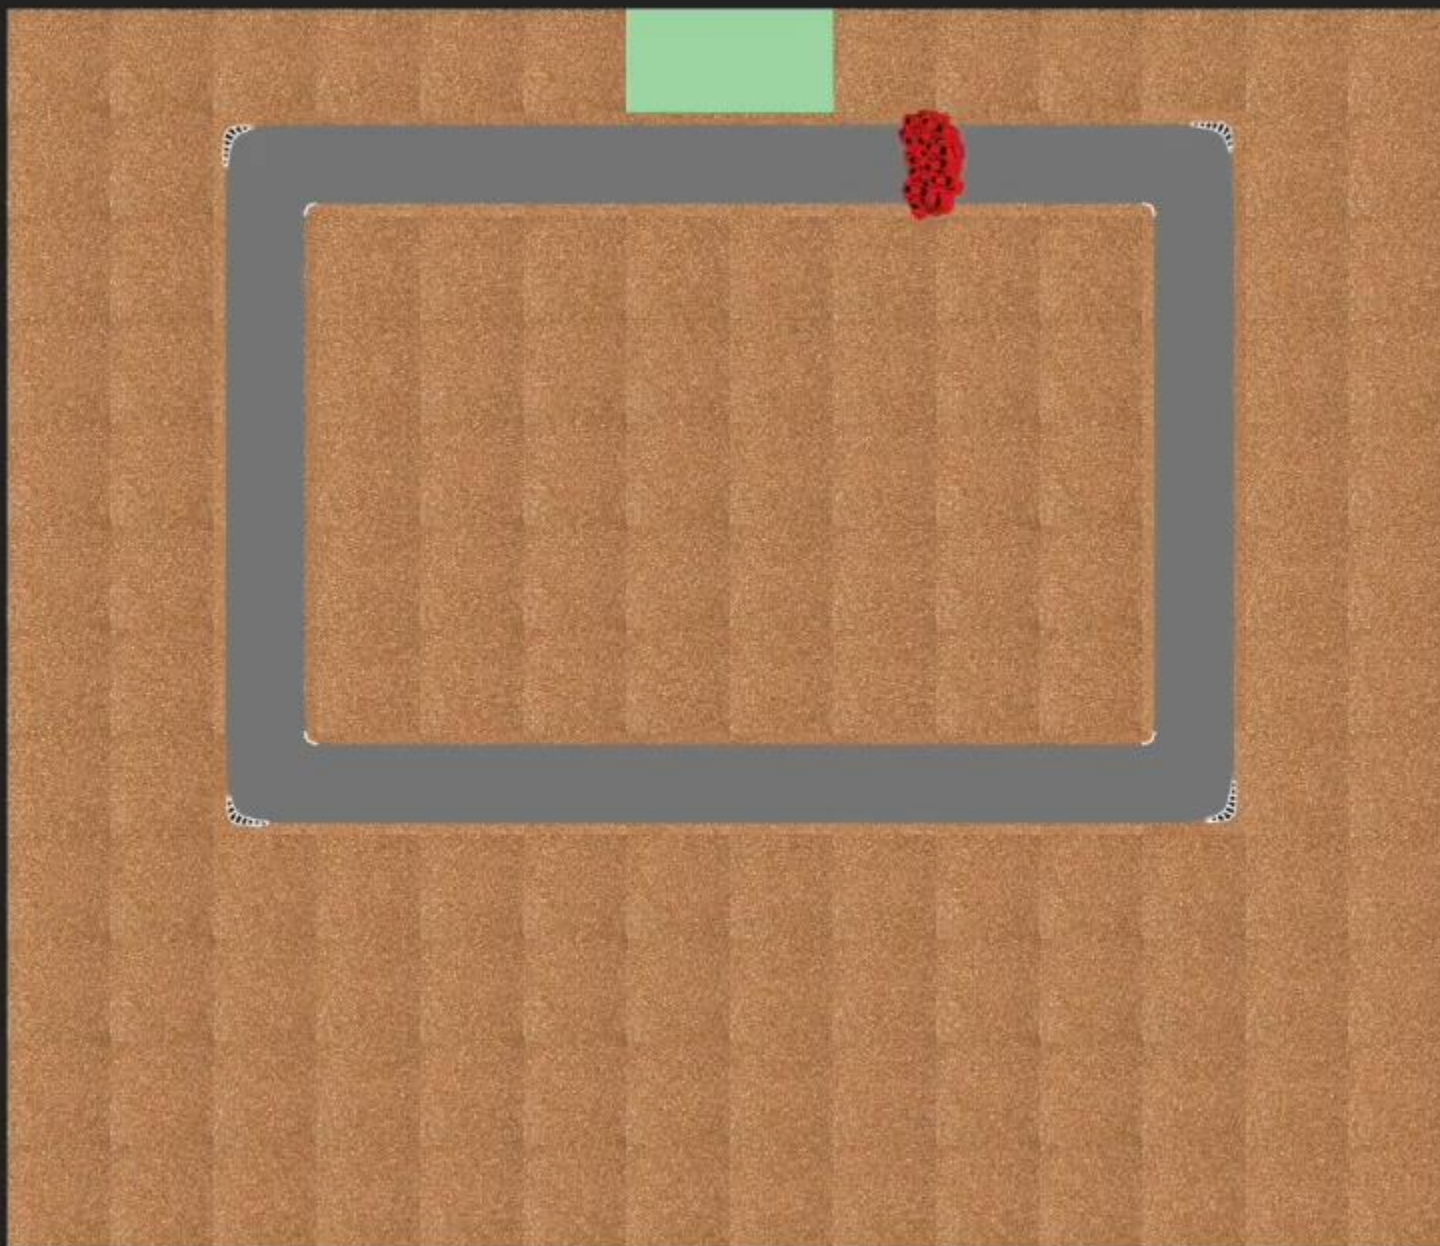
Step-size control

- For updating the mean and covariance matrix between generations



if several updates go to the same/similar direction the step-size is increased

We've seen this before (momentum)



Mutation rate (0.5 - 1.5)

1.0

Number of Elite Cars (0.5 - 1.5)

1.0

Mutation Strength (0.5 - 1.5)

1.0



Update?

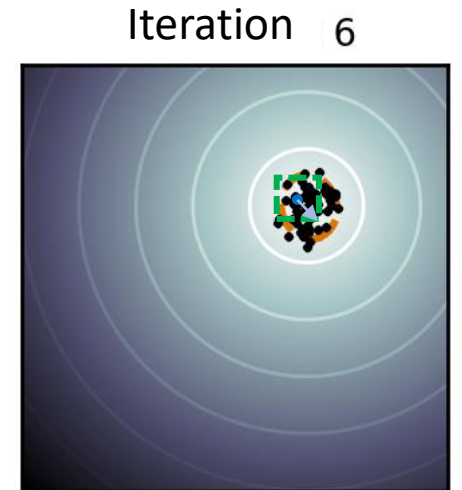
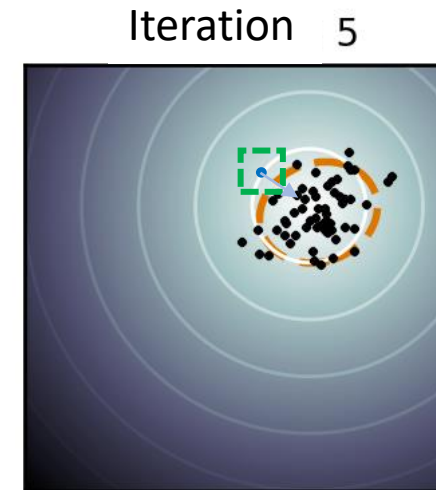
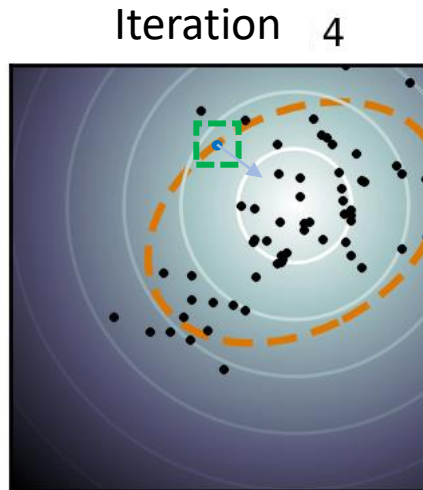
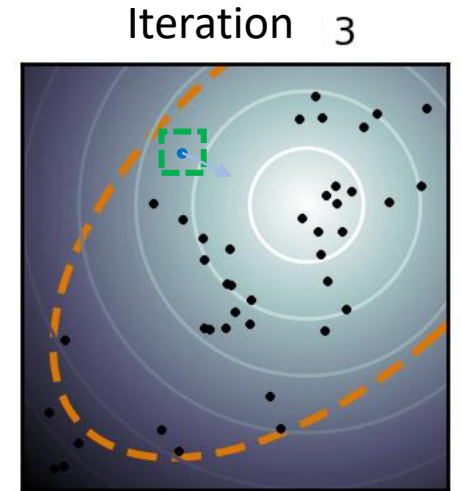
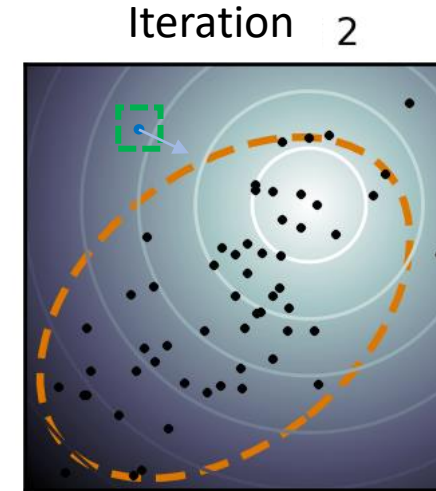
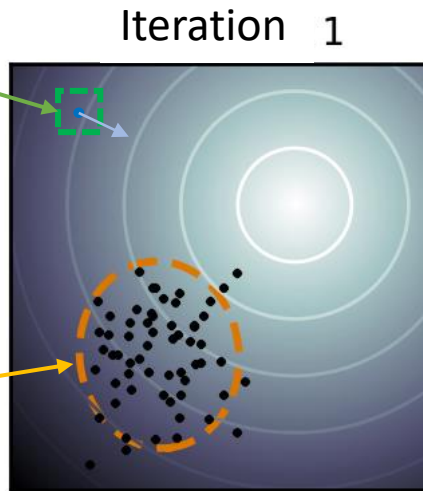
Online parameter tuning

- **Finite difference**

- Less exploration
- Safer
- Harder to escape a local minima

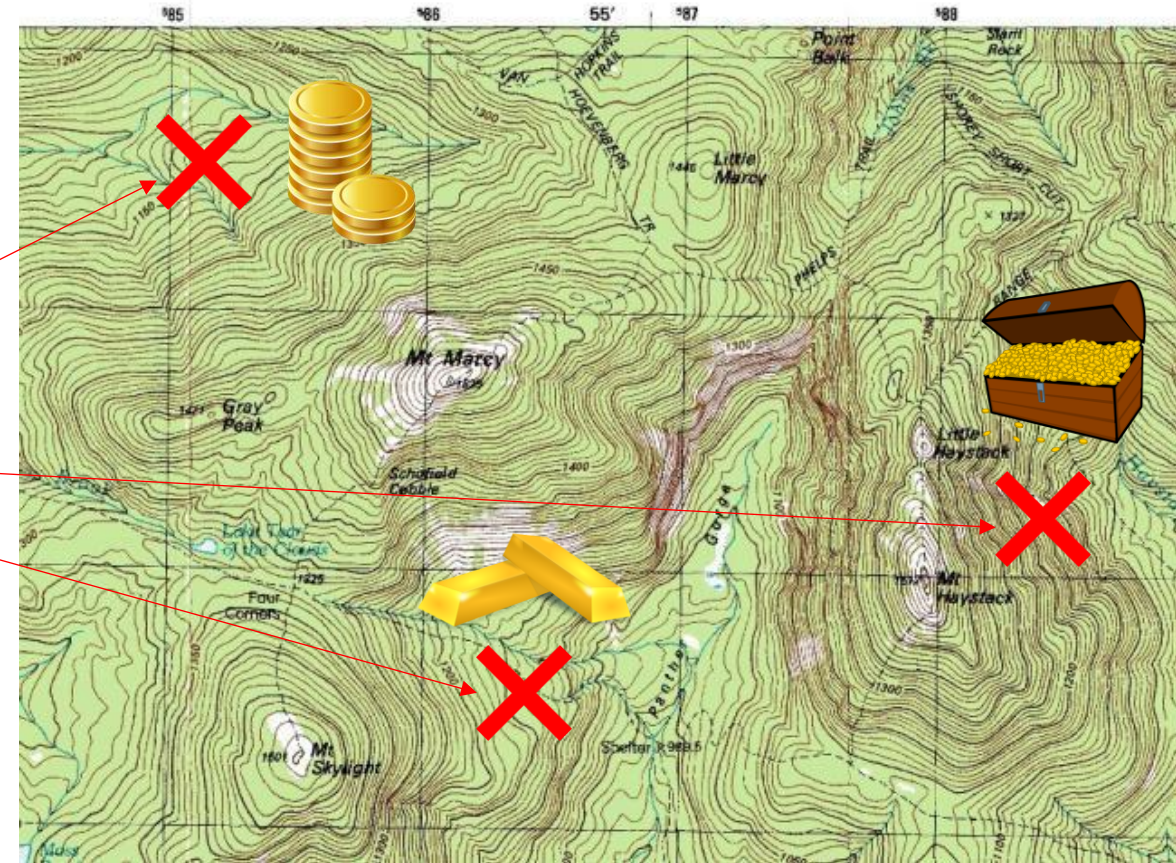
- **CMA-ES**

- More exploration
- Less safe
- Easier to escape a local minima



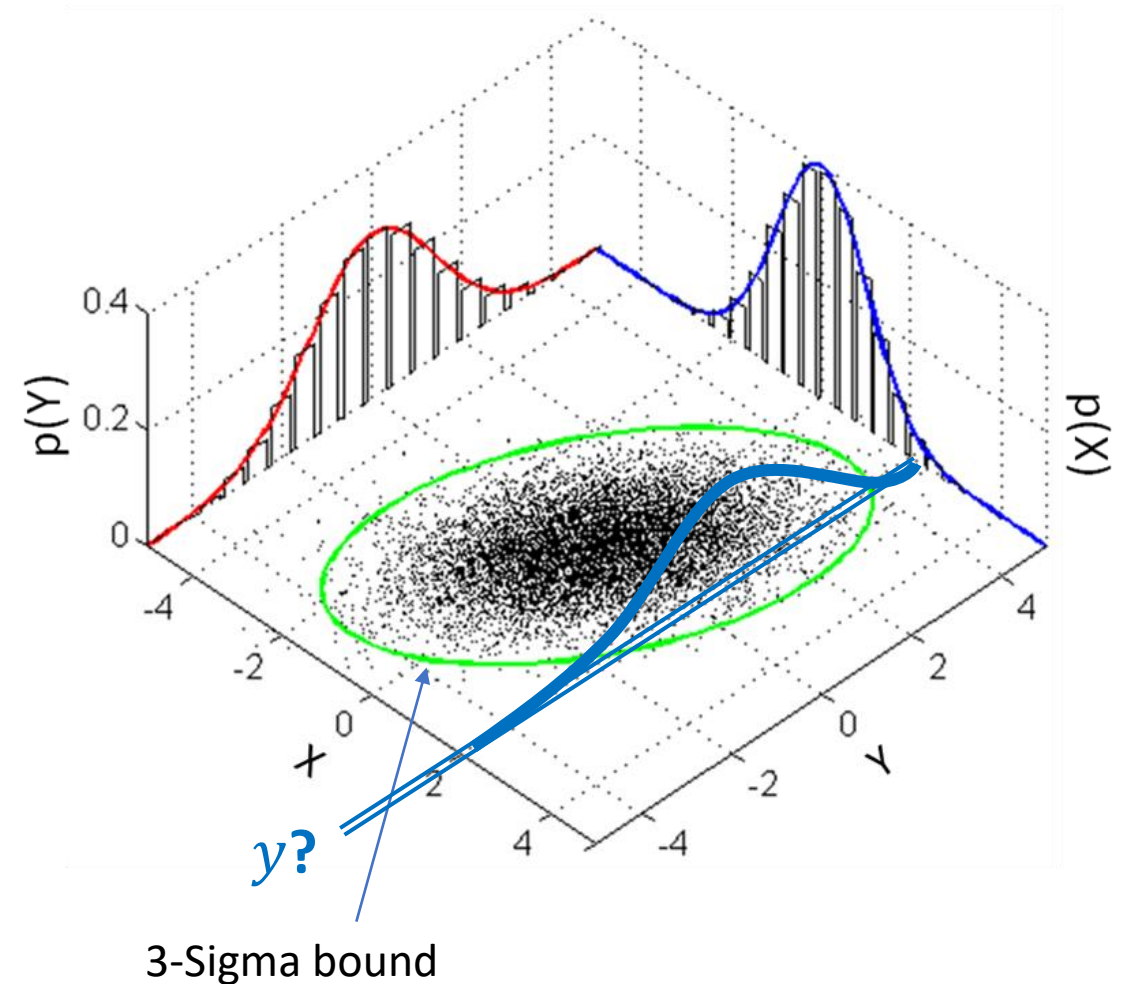
Looking for gold

- Daniel Krige - South African mining engineer, 1960
- Find the point richest in gold
- Sampling the soil is expensive
- Given current samples
 - Where should the next sample be taken from?



Conditional distribution

- Multivariate normal distribution
- $\begin{bmatrix} x \\ y \end{bmatrix} \sim \mathcal{N} \left(\begin{bmatrix} 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 1 & 3/5 \\ 3/5 & 2 \end{bmatrix} \right)$
- What is the distribution over y given $x = 2$?



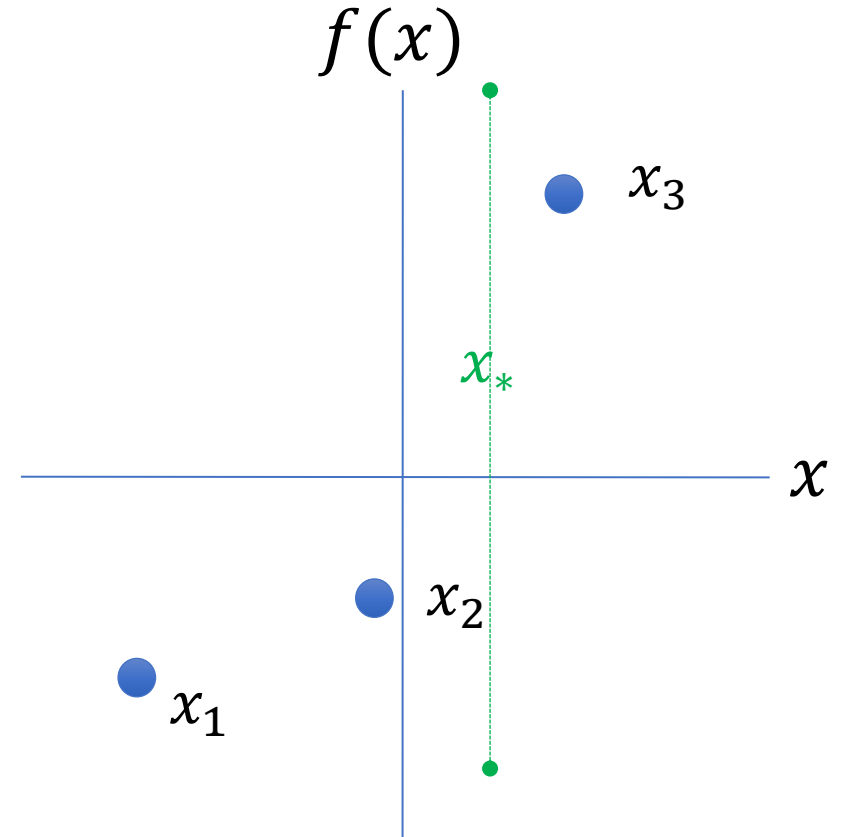
Posterior conditionals of an MVN

- Theorem 4.3.1, “Machine Learning: a Probabilistic Perspective” by Kevin Patrick Murphy
- $\mu_{x|y} = \mu_x + C_{xy}C_{yy}^{-1}(y - \mu_y)$
- $C_{x|y} = C_{xx} - C_{xy}C_{yy}^{-1}C_{yx}$
- $\begin{bmatrix} x \\ y \end{bmatrix} \sim \mathcal{N} \left(\begin{bmatrix} 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 1 & 3/5 \\ 3/5 & 2 \end{bmatrix} \right)$
- $\mu_{x|y} =$
- $C_{x|y} =$

In the example x and y are a single parameter but they can represent high-dimensional vectors (the general case). Hence the matrix notation.

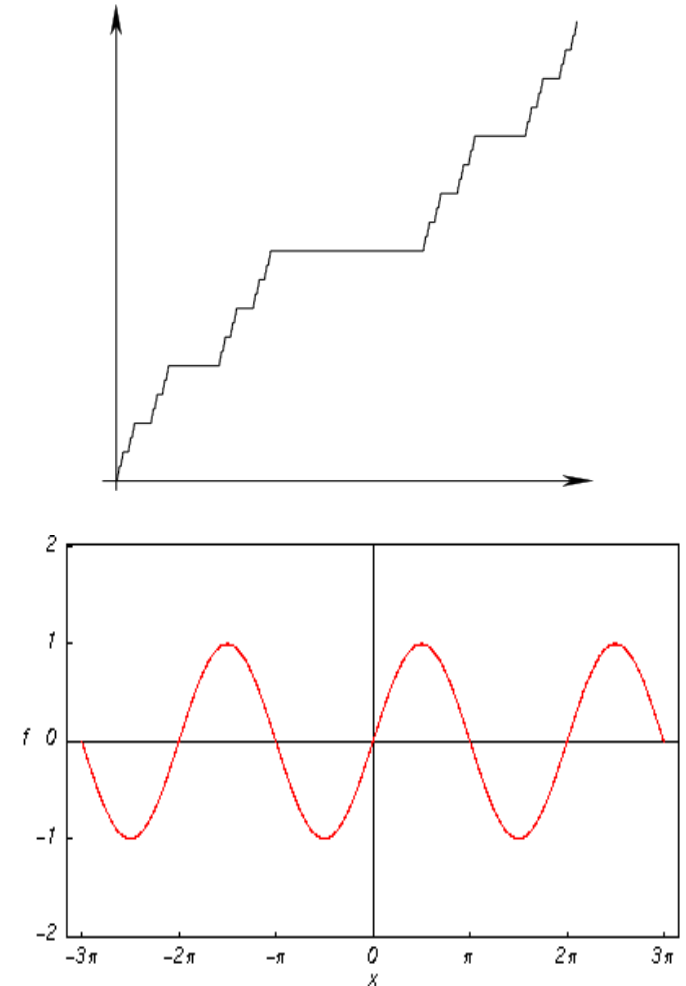
Gaussian Prediction

- Given training data $\{X, f_i\}$ and x_* predict the mean and variance for f_*
- “The key idea is that if x_i and x_j are deemed by the kernel to be similar, then we expect the output of the function at those points to be similar, too”
Machine Learning: a Probabilistic Perspective



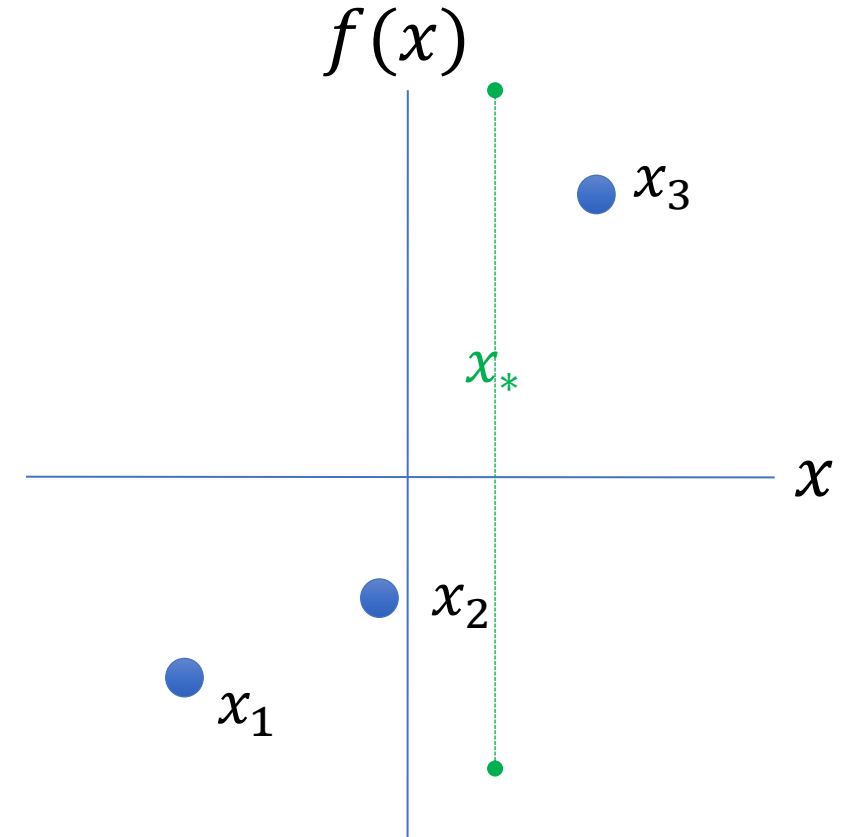
Kernels

- Kernel functions define the co-variance between two points in some (user defined) space
- The chosen space should fit the approximated function/domain
- E.g., The **Squared Exponential Kernel**
- $k(x_1, x_2) = \exp(-\|x_1 - x_2\|^2)$
 - $= 0$, when $\|x_1 - x_2\| \rightarrow \infty$
 - $= 1$, when $\|x_1 - x_2\| \rightarrow 0$



Gaussian Prediction

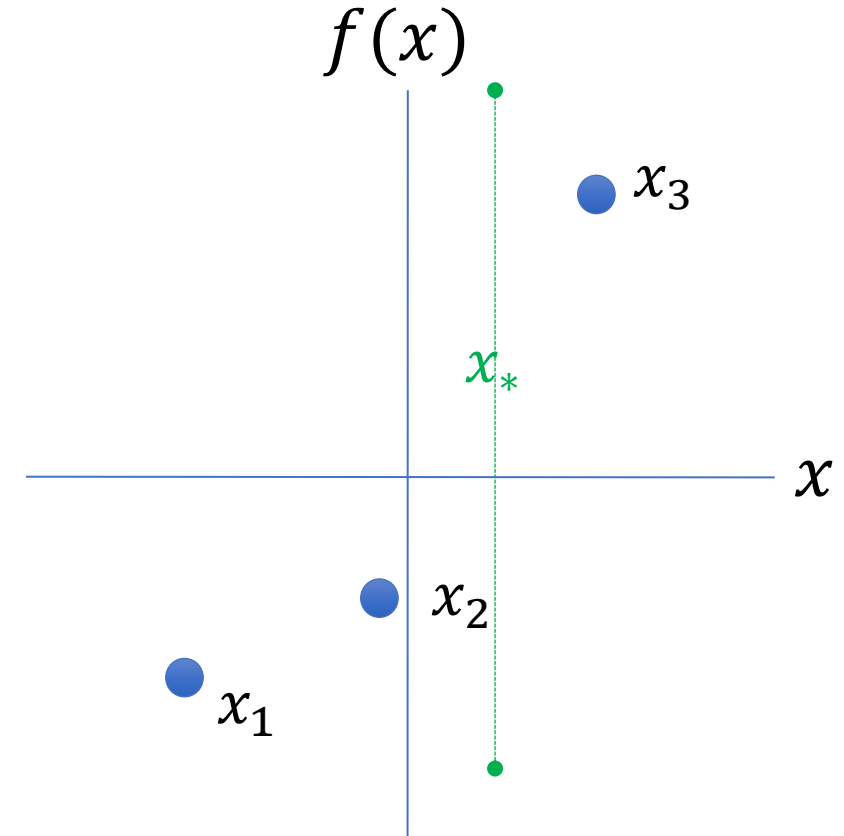
- $k = \begin{bmatrix} 1 & 0.7 & 0.2 \\ 0.7 & 1 & 0.6 \\ 0.2 & 0.6 & 1 \end{bmatrix}$
- Assume: $f \sim \mathcal{N}(\mu, k)$
- Assume: $f_* \sim \mathcal{N}(\mu, k(x_*, x_*))$
 - $k(x_*, x_*) = 1$
 - $\mathcal{N}(\mu, 1)$ is not very helpful
- We assume that f and f_* are jointly Gaussian



Gaussian Process

- $\begin{bmatrix} f \\ f_* \end{bmatrix} \sim \mathcal{N} \left(\mu, \begin{bmatrix} k_{11} & k_{12} & k_{13} & k_{1*} \\ k_{21} & k_{22} & k_{23} & k_{2*} \\ k_{31} & k_{32} & k_{33} & k_{3*} \\ k_{*1} & k_{*2} & k_{*3} & k_{**} \end{bmatrix} \right)$
- $k_{**} = k(x_*, x_*)$
- $k_{1*} = k(x_1, x_*)$
- How can we determine the mean and variance for f_* ?

k k_* k_{**}



Gaussian Process

- Posterior conditionals of a MVN!

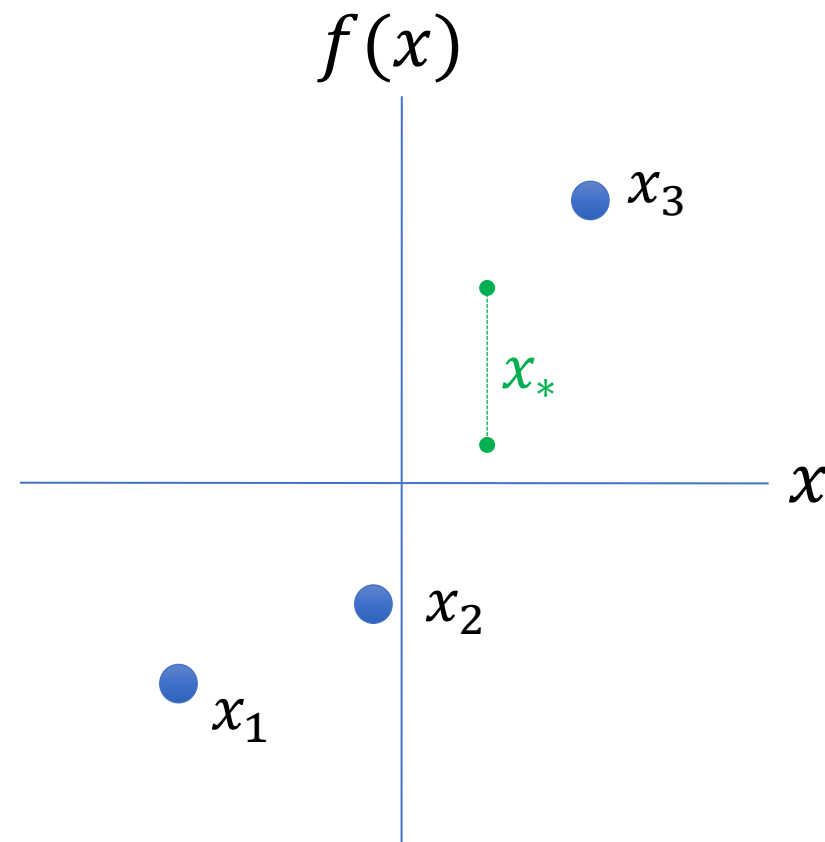
- $\mu_{x|y} = \mu_x + C_{xy}C_{yy}^{-1}(y - \mu_y)$

- $C_{x|y} = C_{xx} - C_{xy}C_{yy}^{-1}C_{yx}$

- $\begin{bmatrix} f \\ f_* \end{bmatrix} \sim \mathcal{N} \left(0, \begin{bmatrix} k & k_* \\ k_*^T & k_{**} \end{bmatrix} \right)$

- $E[f_*] | f = k_*^T k^{-1} f$

- $c_* | f = k_{**} - k_*^T k^{-1} k_*$

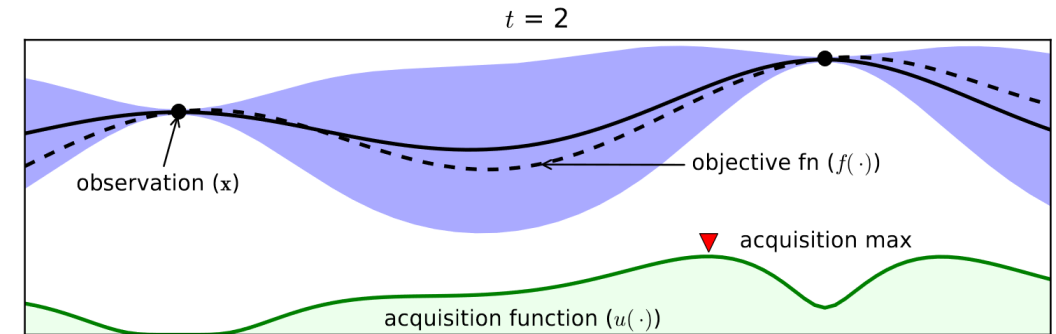


Bayesian optimization

- A useful tool when:
 - f (the objective function) is “expensive to evaluate”
 - Simulate a day’s traffic, car crash outcome, drill a hole, human subjects
 - Bounded number of samples
 - f is continuous
 - f lacks known special structure like concavity or linearity
 - No known first- or second-order derivatives
 - Global optimum is required

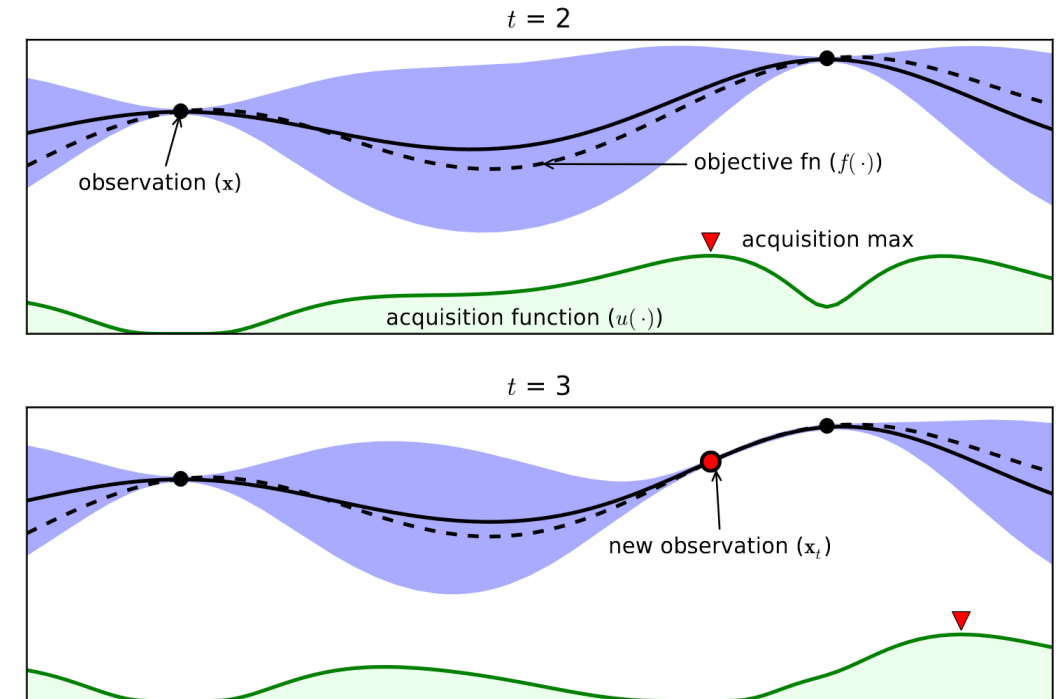
Bayesian optimization

1. For $t = 1$ until n
 2. Get the best sampling candidate: $x_t = \operatorname{argmax}_x u(x | \mathcal{D}_{1:t-1})$
 3. Sample the objective function: $y_t = f(x_t)$
 4. Update the GP according to: $\mathcal{D}_{1:t} = \{\mathcal{D}_{1:t-1}, (x_t, y_t)\}$



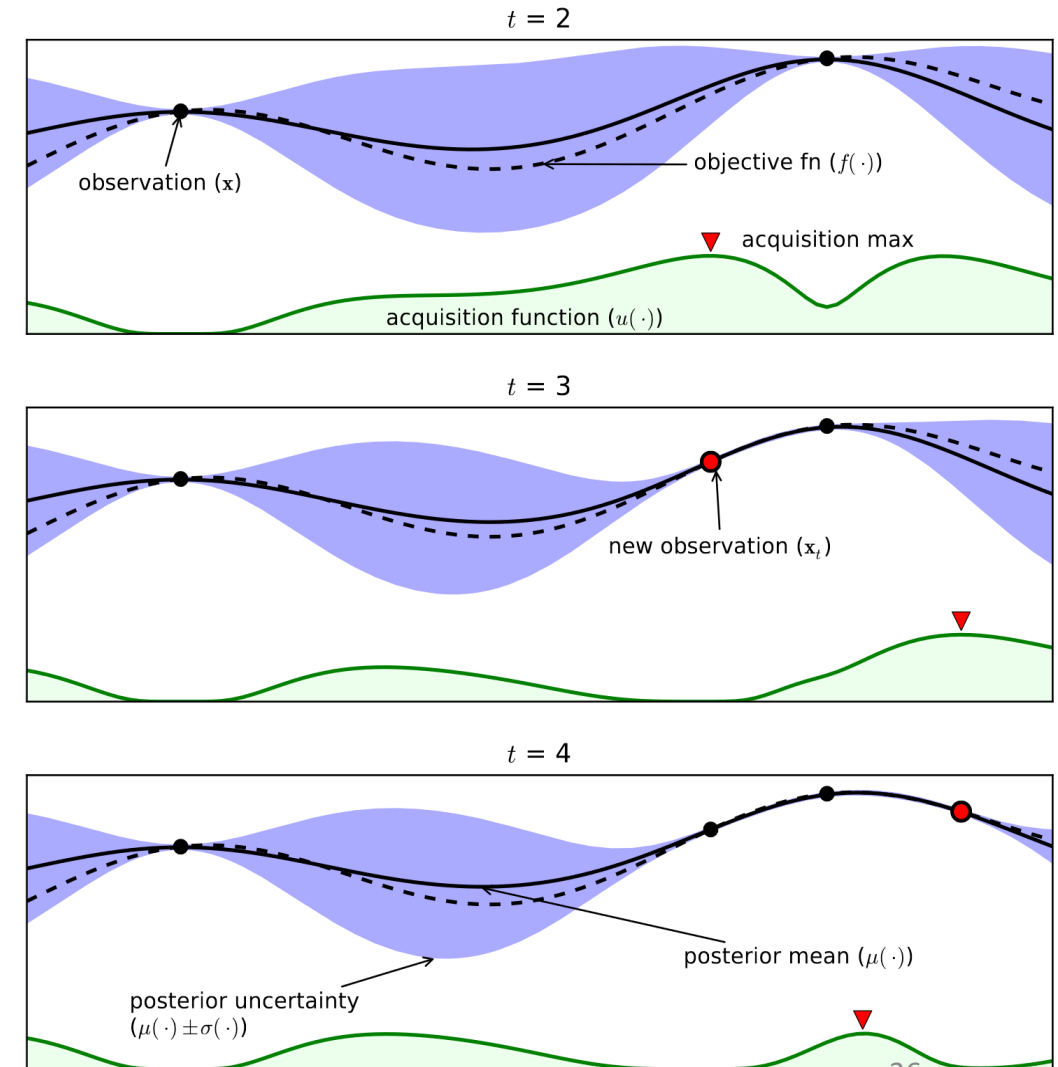
Bayesian optimization

1. For $t = 1$ until n
2. Get the best sampling candidate: $x_t = \operatorname{argmax}_x u(x | \mathcal{D}_{1:t-1})$
3. Sample the objective function: $y_t = f(x_t)$
4. Update the GP according to: $\mathcal{D}_{1:t} = \{\mathcal{D}_{1:t-1}, (x_t, y_t)\}$



Bayesian optimization

1. For $t = 1$ until n
2. Get the best sampling candidate: $x_t = \operatorname{argmax}_x u(x | \mathcal{D}_{1:t-1})$
3. Sample the objective function: $y_t = f(x_t)$
4. Update the GP according to: $\mathcal{D}_{1:t} = \{\mathcal{D}_{1:t-1}, (x_t, y_t)\}$

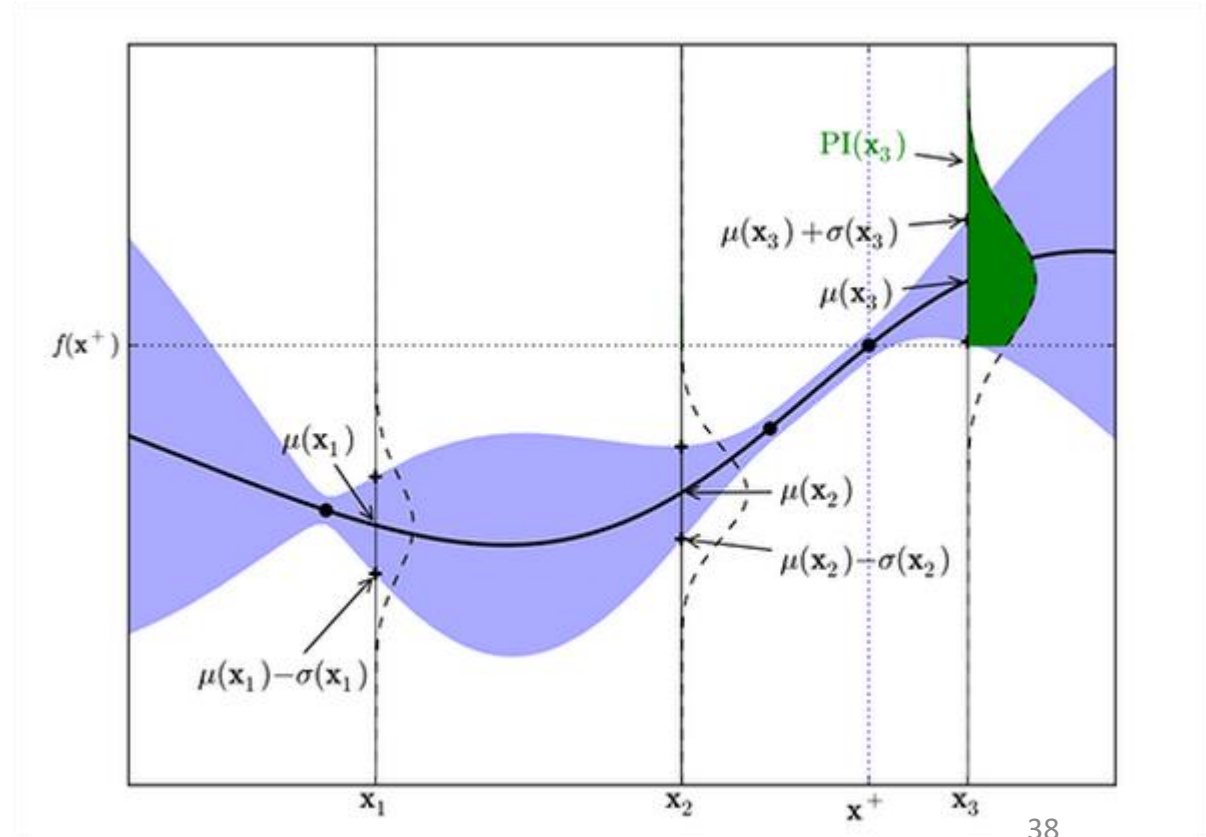


Acquisition function

- **Exploit:** prefer the points with high expected mean value
- **Explore:** prefer the points with high variance
- The acquisition function balances between the two
 1. Probability of Improvement
 2. Expected Improvement
 3. Gaussian Process Upper Confidence Bound (GP-UCB)
 4. Thompson sampling

Probability of Improvement [Kushner, 1964]

- $PI(x_*) = P(f(x_*) > \mu^+ + \xi) = \Phi\left(\frac{\mu_* - \mu^+ - \xi}{\text{Var}_*}\right)$
 - Φ is the normal CDF
- The probability that x_* is better than the best known point (μ^+)
- ξ is required so to bias against previously observed points
- PI is very useful if the maximal f value is known a priori



Expected Improvement [Mockus 1978]

- $x_* = \operatorname{argmax}_x \mathbb{E}(\max\{0, f_{n+1}(x) - f^{max}\} | \mathcal{D}_n)$
 - $f^{max} = \max(\mu^+ + \xi)$
- If X is a random variable whose cumulative distribution function admits a density $f(x)$ then: $\mathbb{E}(x) = \int x f(x) dx$
- $EI(x_*) = \begin{cases} (\mu_* - \mu^+ - \xi)\Phi(Z) + Var_*\phi(Z) & Var_* > 0 \\ 0 & Var_* = 0 \end{cases}$
 - $Z = \frac{\mu_* - \mu^+ - \xi}{\sqrt{Var_*}}$
 - ϕ is the normal PDF

GP-UCB [Srinivas et al. 2010]

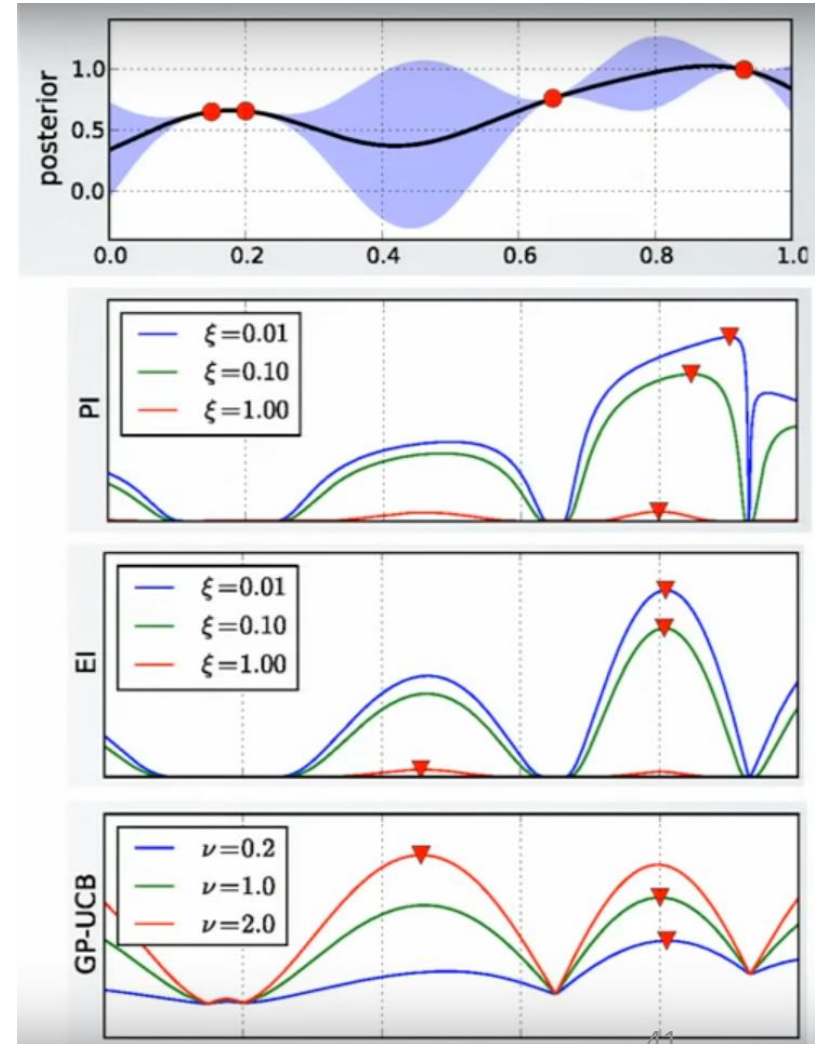
- $GPUCB(x_*) = \mu_* + \sqrt{v\beta_t}Var_*$
- Linear combination of the mean and variance
- Setting $v = 1$ and $\beta_t = 2 \log \left(\frac{t^{d/2+2}\pi^2}{3\delta} \right)$ leads to no regret:

$$\lim_{T \rightarrow \infty} \frac{R_T}{T} = 0$$

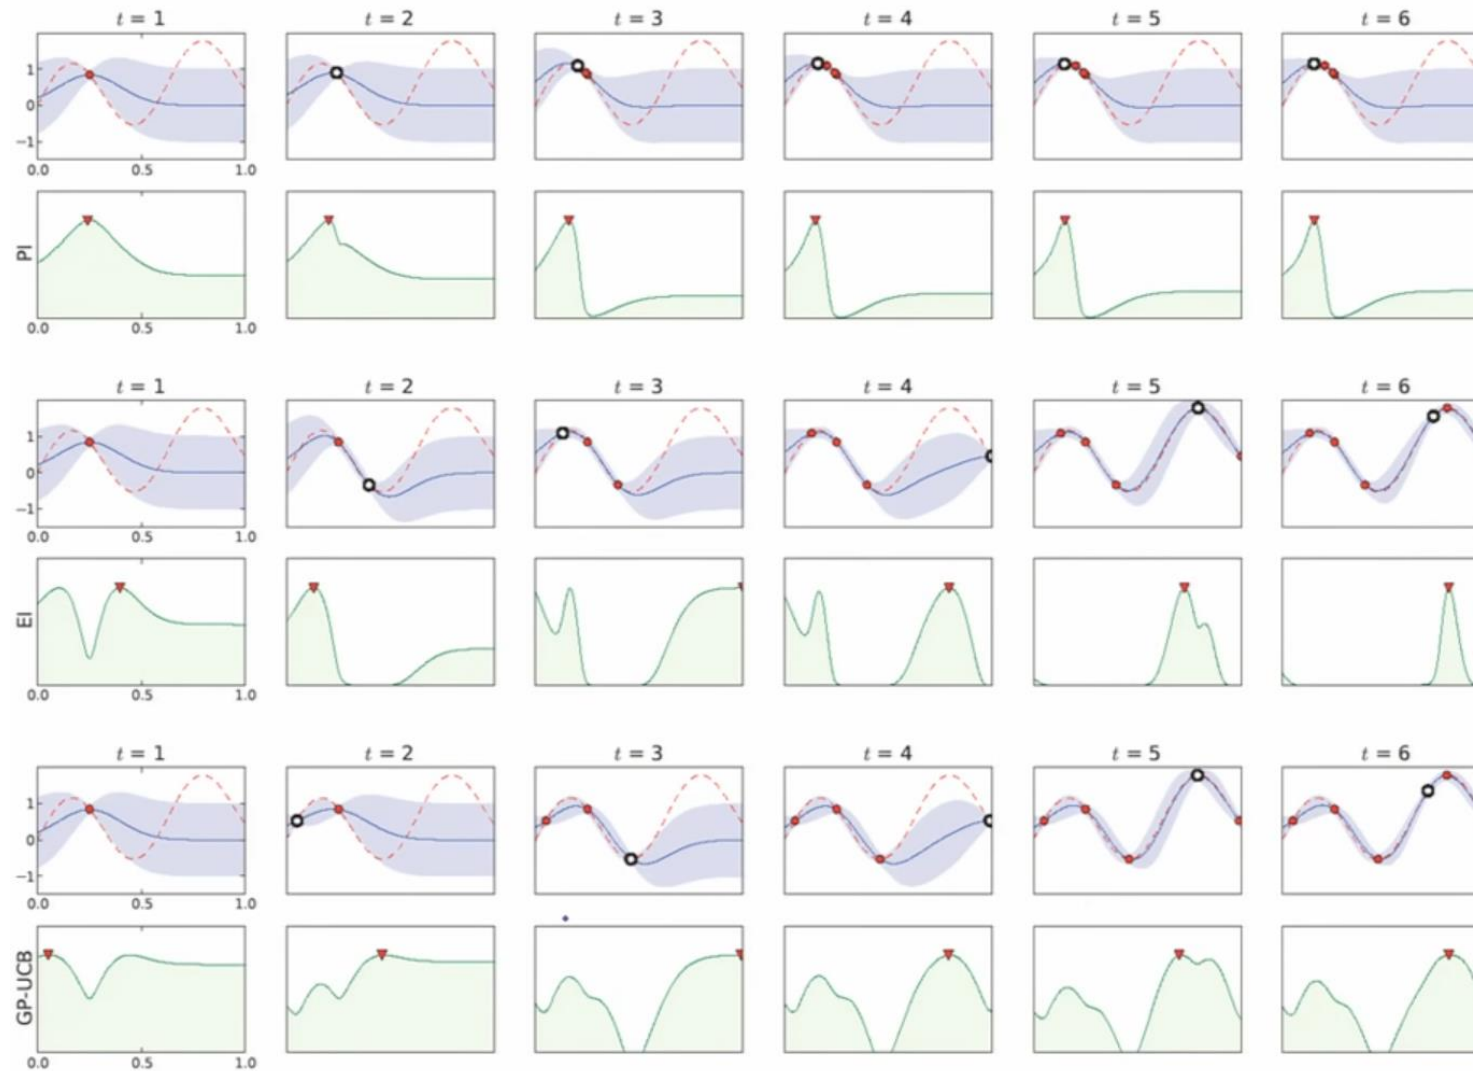
- $R_T = r(x_1) + \dots + r(x_T)$
- $r(x) = f(x^*) - f(x)$

Acquisition functions

- **PI**: Probability improvement
- **EI**: Expected improvement
- **GP-UCB**: Gaussian process upper confidence bound

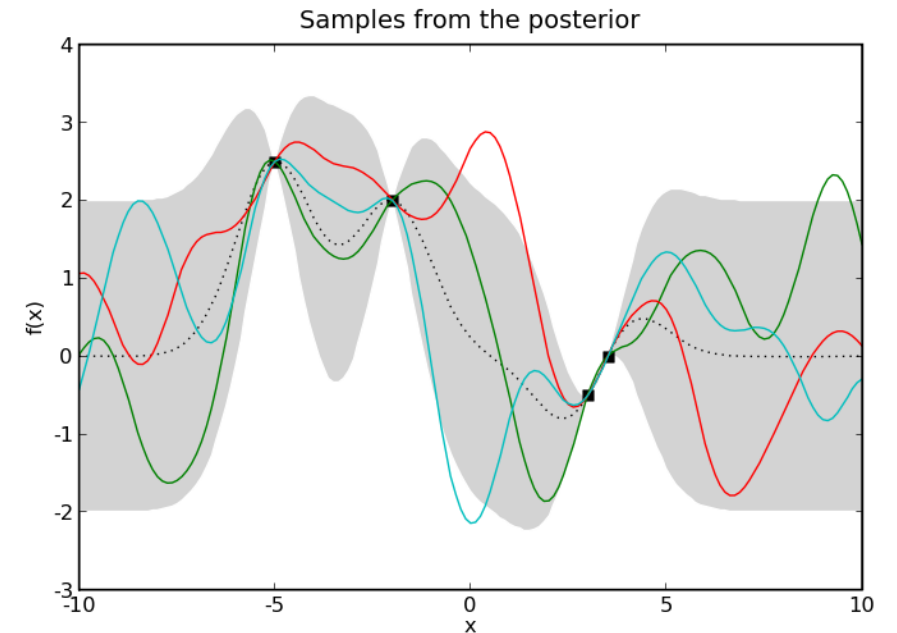


Acquisition functions



Thompson Sampling

- Sample a function from the Gaussian Process
- The function is sampled as a set of f_* values from the distribution
$$\begin{bmatrix} f \\ f_* \end{bmatrix} \sim \mathcal{N} \left(\begin{bmatrix} \mu \\ \mu_* \end{bmatrix}, \begin{bmatrix} k & k_* \\ k_*^T & k_{**} \end{bmatrix} \right) \text{ given } f$$
- Chosen point is the optimum of the sampled function



Extra reading

- https://www.youtube.com/watch?v=SQtOI9jsrJ0&list=PLCJPYIcPhgPAiiBjjVxi5St_IC_cXHfCK&index=9
- <https://www.youtube.com/watch?v=4vGiHC35j9s>

What next?

- **Lecture:** Curriculum Learning
- **Assignments:**
 - DDPG, by No. 25, EoD
- **Quiz (on Canvas):**
 - Imitation Learning, by Nov 25, EoD
- **Project:**
 - Final Report, by Dec. 2, EoD